



Target
Insight

Target Sales Analytics Capstone

Data Analytics with Advanced SQL | Power BI | MySQL |
DAX | Visualization | python | Machine-Learning

www.TargetSalesInsights.com



By: K K Sahani

Overview

- Introduction
- Dataset Description
- Phase 1: SQL Analysis
- Phase 2: Data Visualization
- Phase 3: Predictive Analytics
- Reference






Introduction

This project focuses on analyzing sales, customer behavior, and operational performance using data visualization and analytics. By leveraging datasets such as orders, products, customers, and regions, the goal is to uncover actionable insights that drive informed decision-making. Key areas of analysis include sales trends, customer loyalty, revenue by category, regional performance, and operational metrics like delivery times and return rates.

The project employs Power BI for interactive dashboards and DAX queries for advanced calculations, enabling a clear understanding of business performance and opportunities for optimization. It is designed to assist stakeholders in identifying growth areas, enhancing customer experiences, and improving operational efficiency.





Phase 1: SQL Analysis – Advanced Queries

Introduction:

This phase focuses on leveraging advanced SQL queries to analyze sales, customer behavior, and operational performance. Key tasks include identifying top-selling products, loyal customers, and sales trends, as well as evaluating delivery efficiency and cancellations. Advanced techniques like window functions and time-based metrics (MTD, YTD) are used to extract actionable insights for informed decision-making.




About MySQL

- **Powerful and Open-Source:** A widely used relational database management system (RDBMS).
- **Data Storage and Management:** Used for storing and managing structured data.
- **SQL Querying:** Provides robust tools for querying, analysing, and manipulating data through SQL (Structured Query Language).
- **Scalability:** Can handle large datasets and grow with business needs.
- **Reliability:** Known for its stable performance and consistency.
- **Efficiency:** Optimized for fast data processing and complex business queries.

1. Sales Performance Analysis:

1. Write a query to calculate total sales revenue per-category, sub-category, and region.

```
SELECT
  p.product_category,
  c.customer_state AS region,
  SUM(pay.payment_value) AS total_sales_revenue
FROM
  ecommerce_sales.order_items oi
JOIN
  ecommerce_sales.products p ON oi.product_id = p.product_id
JOIN
  ecommerce_sales.orders o ON oi.order_id = o.order_id
JOIN
  ecommerce_sales.customers c ON o.customer_id = c.customer_id
JOIN
  ecommerce_sales.payments pay ON o.order_id = pay.order_id
WHERE
  o.order_status = 'delivered'
GROUP BY
  p.product_category,
  c.customer_state
LIMIT 100;
```




| product_category | customer_state | Total_sales_revenue |
|------------------|----------------|---------------------|
| sport leisure | SP | 3750448.6950540543 |
| bed table bath | BA | 249495.2992630005 |
| bed table bath | RJ | 1442517.9023265839 |
| HEALTH BEAUTY | RJ | 1402363.499970436 |
| bed table bath | SP | 4722380.698928833 |
| automotive | GO | 113092.59939193726 |

1. Sales Performance Analysis:

2. Identify the top 5 best-selling products by both sales revenue and quantity sold.

-- Top 5 Best-Selling Products by Sales Revenue


```
SELECT
  p.product_id,
  p.product_category AS category,
  SUM(oi.price * oi.order_item_id) AS total_revenue
FROM
  order_items oi
JOIN
  products p
ON
  oi.product_id = p.product_id
GROUP BY
  p.product_id, p.product_category
ORDER BY
  total_revenue DESC
LIMIT 5;
```



| product_id | category | total_revenue |
|----------------------------------|----------------------|-------------------|
| bb50f2e236e5eea0100680137654686c | HEALTH BEAUTY | 6343650 |
| 5769ef0a239114ac3a854af00df129e4 | fixed telephony | 5443200 |
| 6cdd53843498f92890544667809f1595 | HEALTH BEAUTY | 5180183.931884766 |
| d1c427060a0f73f6b889a5c7c61f2ac4 | computer accessories | 4584635.148696899 |
| d6160fb7873f184099d9bc95e30376af | PCs | 4400940.596923828 |

-- Top 5 Best-Selling Products by Quantity Sold

```
SELECT
  p.product_id,
  p.product_category AS category,
  SUM(oi.order_item_id) AS total_quantity_sold
FROM
  order_items oi
JOIN
  products p
ON
  oi.product_id = p.product_id
GROUP BY
  p.product_id, p.product_category
ORDER BY
  total_quantity_sold DESC
LIMIT 5;
```



| product_id | category | total_quantity_sold |
|----------------------------------|----------------------|---------------------|
| 422879e10f46682990de24d770e7f83d | Garden tools | 71370 |
| aca2eb7d00ea1a7b8ebd4e68314663af | Furniture Decoration | 57600 |
| 368c6c730842d78016ad823897a372db | Garden tools | 49590 |

2. Customer Insights:

1. Find the most loyal customers by calculating their purchase frequency and total spend.

- Most Loyal Customers (Purchase Frequency and Total Spend)

```
SELECT
  c.customer_id,
  c.customer_unique_id,
  COUNT(o.order_id) AS purchase_frequency,
  SUM(oi.price * oi.order_item_id) AS total_spend
FROM
  customers c
JOIN
  orders o
ON
  c.customer_id = o.customer_id
JOIN
  order_items oi
ON
  o.order_id = oi.order_id
GROUP BY
  c.customer_id, c.customer_unique_id
ORDER BY
  purchase_frequency DESC, total_spend DESC
LIMIT 10; -- Top 10 most loyal customers
```




| customer_id | purchase_freq | total_spend |
|----------------------------------|---------------|--------------------|
| fc3d1daec319d62d49bfb5e1f83123e9 | 210.0000 | 318.0000114440918 |
| be1b70680b9f9694d8c70f41fa3dc92b | 200.0000 | 20000 |
| bd5d39761aa56689a265d95d8d32b8be | 200.0000 | 19739.999389648438 |
| 10de381f8a8d23fff822753305f71cae | 150.0000 | 9823.49967956543 |
| adb32467ecc74b53576d9d13a5a55891 | 150.0000 | 7650 |
| d5f2b3f597c7ccafb5cac0bcc3d6024 | 140.0000 | 8260 |

2. Customer Insights:

2. Identify customers with the highest average order value (AOV).

-- Query 2: Customers with the Highest AOV

```
SELECT
  c.customer_id,
  c.customer_unique_id,
  SUM(oi.price * oi.order_item_id) / COUNT(DISTINCT o.order_id)
AS average_order_value,
  SUM(oi.price * oi.order_item_id) AS total_spend,
  COUNT(DISTINCT o.order_id) AS total_orders
FROM
  customers c
JOIN
  orders o
ON
  c.customer_id = o.customer_id
JOIN
  order_items oi
ON
  o.order_id = oi.order_id
GROUP BY
  c.customer_id, c.customer_unique_id
ORDER BY
  average_order_value DESC
LIMIT 10; -- Top 10 customers with the highest AOV
```




| customer_id | AOV |
|----------------------------------|-------------------|
| 1617b1357756262bfa56ab541c47bc16 | 134400 |
| ec5b2ba62e574342386871631fafd3fc | 71600 |
| c6e2731c5b391845f6800c97401a43a9 | 67350 |
| f48d464a0baaea338cb25f816991ab1f | 67290 |
| 3fd6777bbce08a352fddd04e4a7cc8f6 | 64990 |
| 05455dfa7cd02f13d132aa7a6a9729c6 | 59345.99853515625 |
| df55c14d1476a9a3467f131269c2477f | 47990 |

3. Operational Efficiency:

1. Analyze delivery performance by calculating the average delivery time by region.

• -- Average Delivery Time by Region

```
WITH delivery_time AS (  
  SELECT  
    o.customer_id,  
    DATEDIFF(o.order_delivered_customer_date,  
o.order_purchase_timestamp) AS delivery_time  
  FROM  
    orders o  
  WHERE  
    o.order_status = 'delivered'  
    AND o.order_delivered_customer_date IS NOT NULL  
)  
SELECT  
  c.customer_state AS region,  
  AVG(d.delivery_time) AS avg_delivery_time  
FROM  
  delivery_time d  
JOIN  
  customers c  
  ON d.customer_id = c.customer_id  
GROUP BY  
  c.customer_state  
ORDER BY  
  avg_delivery_time ASC;
```




| Result Grid | | Filter Rows: |
|-------------|--------|-------------------|
| | region | avg_delivery_time |
| ▶ | SP | 8.7006 |
| | PR | 11.9380 |
| | MG | 11.9450 |

3. Operational Efficiency:

2..Identify regions or products with the highest canceled rates.

```
--      Regions with the Highest Canceled Rates
SELECT
  c.customer_state AS region,
  COUNT(CASE WHEN o.order_status = 'canceled'
THEN 1 END) * 100.0 / COUNT(*) AS
cancellation_rate
FROM
  orders o
JOIN
  customers c
  ON o.customer_id = c.customer_id
GROUP BY
  c.customer_state
ORDER BY
  cancellation_rate DESC
LIMIT 10;
```




| | region | cancellation_rate |
|---|--------|-------------------|
| ▶ | RR | 2.17391 |
| | RO | 1.18577 |
| | PI | 0.80808 |
| | SP | 0.78331 |
| | RJ | 0.66916 |
| | GO | 0.64356 |
| | MG | 0.55006 |
| | MA | 0.53548 |
| | CE | 0.52395 |
| | SC | 0.52241 |

4. Date and Time Analytics:

1. Write a query to find the monthly sales trend for the last two years.

```
-- Query for Monthly Sales Trend (Last Two Years)
SELECT
  DATE_FORMAT(o.order_purchase_timestamp,
'%Y-%m') AS month,
  SUM(p.payment_value) AS total_sales
FROM
  orders o
JOIN
  payments p
  ON o.order_id = p.order_id
WHERE
  o.order_status = 'delivered'
GROUP BY
  DATE_FORMAT(o.order_purchase_timestamp,
'%Y-%m')
ORDER BY
  month ASC;
```




| Result Grid | | | Filter Rows: |
|-------------|---------|--------------------|--------------|
| | month | total_sales | |
| ▶ | 2016-10 | 5867405.463219166 | |
| | 2016-12 | 2472.120105743408 | |
| | 2017-01 | 16070754.427737594 | |
| | 2017-02 | 34183629.826975465 | |
| | 2017-03 | 52210543.0761053 | |
| | 2017-04 | 49259974.63225865 | |
| | 2017-05 | 71450407.99397416 | |
| | 2017-06 | 61768425.64433563 | |
| | 2017-07 | 71366895.12379843 | |
| | 2017-08 | 81396076.83708212 | |
| Result 1 | | | × |

4. Date and Time Analytics:

2. Analyze the seasonality of sales to identify peak months.

-- Query for Seasonality Analysis (Peak Months)

```
SELECT
  YEAR(order_purchase_timestamp) AS year,
  MONTH(order_purchase_timestamp) AS month,
  SUM(p.payment_value) AS total_sales
FROM
  orders o
JOIN
  payments p
  ON o.order_id = p.order_id
WHERE
  o.order_status = 'delivered'
GROUP BY
  YEAR(order_purchase_timestamp),
  MONTH(order_purchase_timestamp)
ORDER BY
  year, month;
```




| | year | month | total_sales |
|--|------|-------|--------------------|
| | 2016 | 10 | 5867405.463219166 |
| | 2016 | 12 | 2472.120105743408 |
| | 2017 | 1 | 16070754.427737594 |
| | 2017 | 2 | 34183629.826975465 |
| | 2017 | 3 | 52210543.0761053 |
| | 2017 | 4 | 49259974.63225865 |
| | 2017 | 5 | 71450407.99397416 |
| | 2017 | 6 | 61768425.64433563 |
| | 2017 | 7 | 71366895.12379843 |
| | 2017 | 8 | 81396076.83708212 |

result 1 x

5. Advanced SQL Queries:

1. Use window functions to rank products based on their sales within each category.

```
-- -- 1. Rank Products Based on Sales Within Each Category using window function
SELECT
  p.product_category,
  oi.product_id,
  SUM(oi.price) AS total_sales,
  RANK() OVER (PARTITION BY p.product_category ORDER
BY SUM(oi.price) DESC) AS sales_rank
FROM
  order_items oi
JOIN
  products p
  ON oi.product_id = p.product_id
GROUP BY
  p.product_category, oi.product_id
ORDER BY
  p.product_category, sales_rank;
```



| | product_category | product_id | total_sales | sales_rank |
|---|------------------|----------------------------------|--------------------|------------|
| ▶ | NULL | 5a848e4ab52fd5445cdc07aab1c40e48 | 2180612.662124634 | 1 |
| | NULL | eed5cbd74fac3bd79b7c7ec95fa7507d | 895050 | 2 |
| | NULL | b1d207586fca400a2370d50a9ba1da98 | 643680 | 3 |
| | NULL | 76d1a1a9d21ab677a61c3ae34b1b352f | 514137.6013183594 | 4 |
| | NULL | ad88641611c35ebd59ecda07a9f17099 | 406379.71115112305 | 5 |
| | NULL | 3b60d513e90300a4e9833e5cda1f1d61 | 395399.71115112305 | 6 |
| | NULL | 4c50dcc50f1512f46096d6ef0142c4a9 | 358200 | 7 |
| | NULL | 17823ffd2de8234f0e885a71109613a4 | 267290.09170532227 | 8 |
| | NULL | ... | ... | ... |

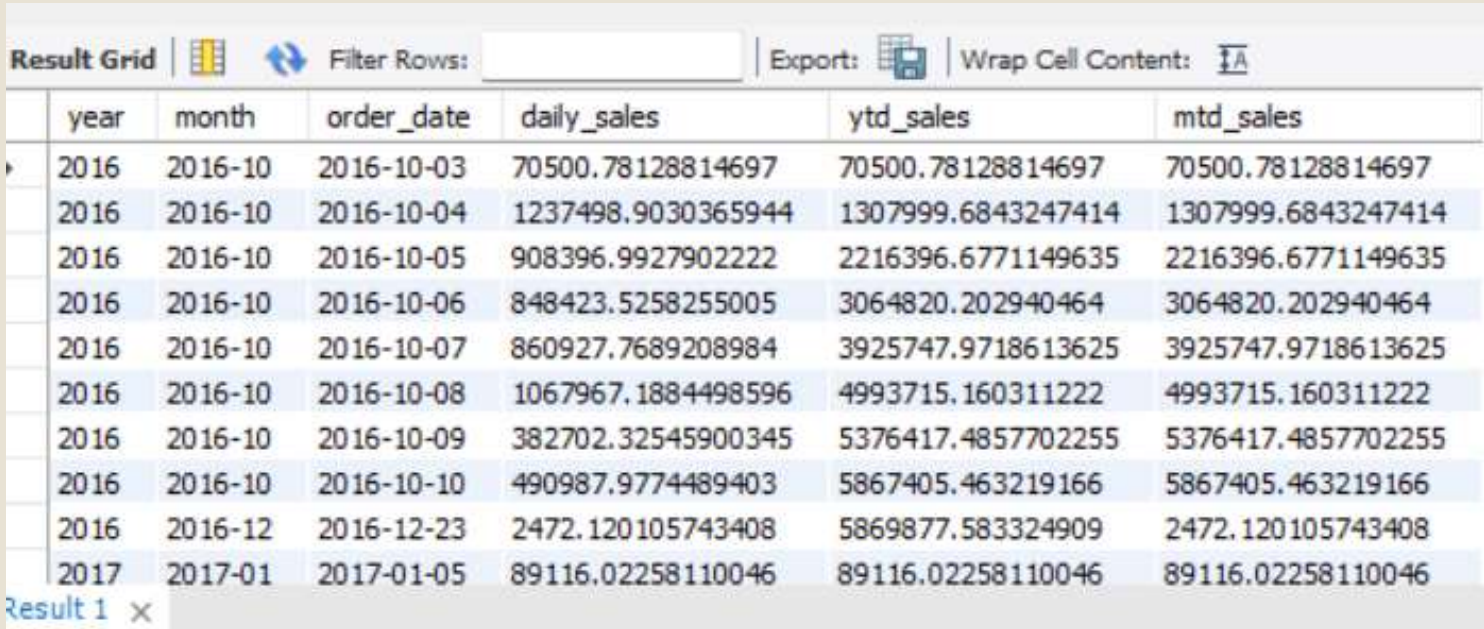
Result 1 x

5. Advanced SQL Queries:

2. Calculate month-to-date (MTD) and year-to-date (YTD) sales metrics.

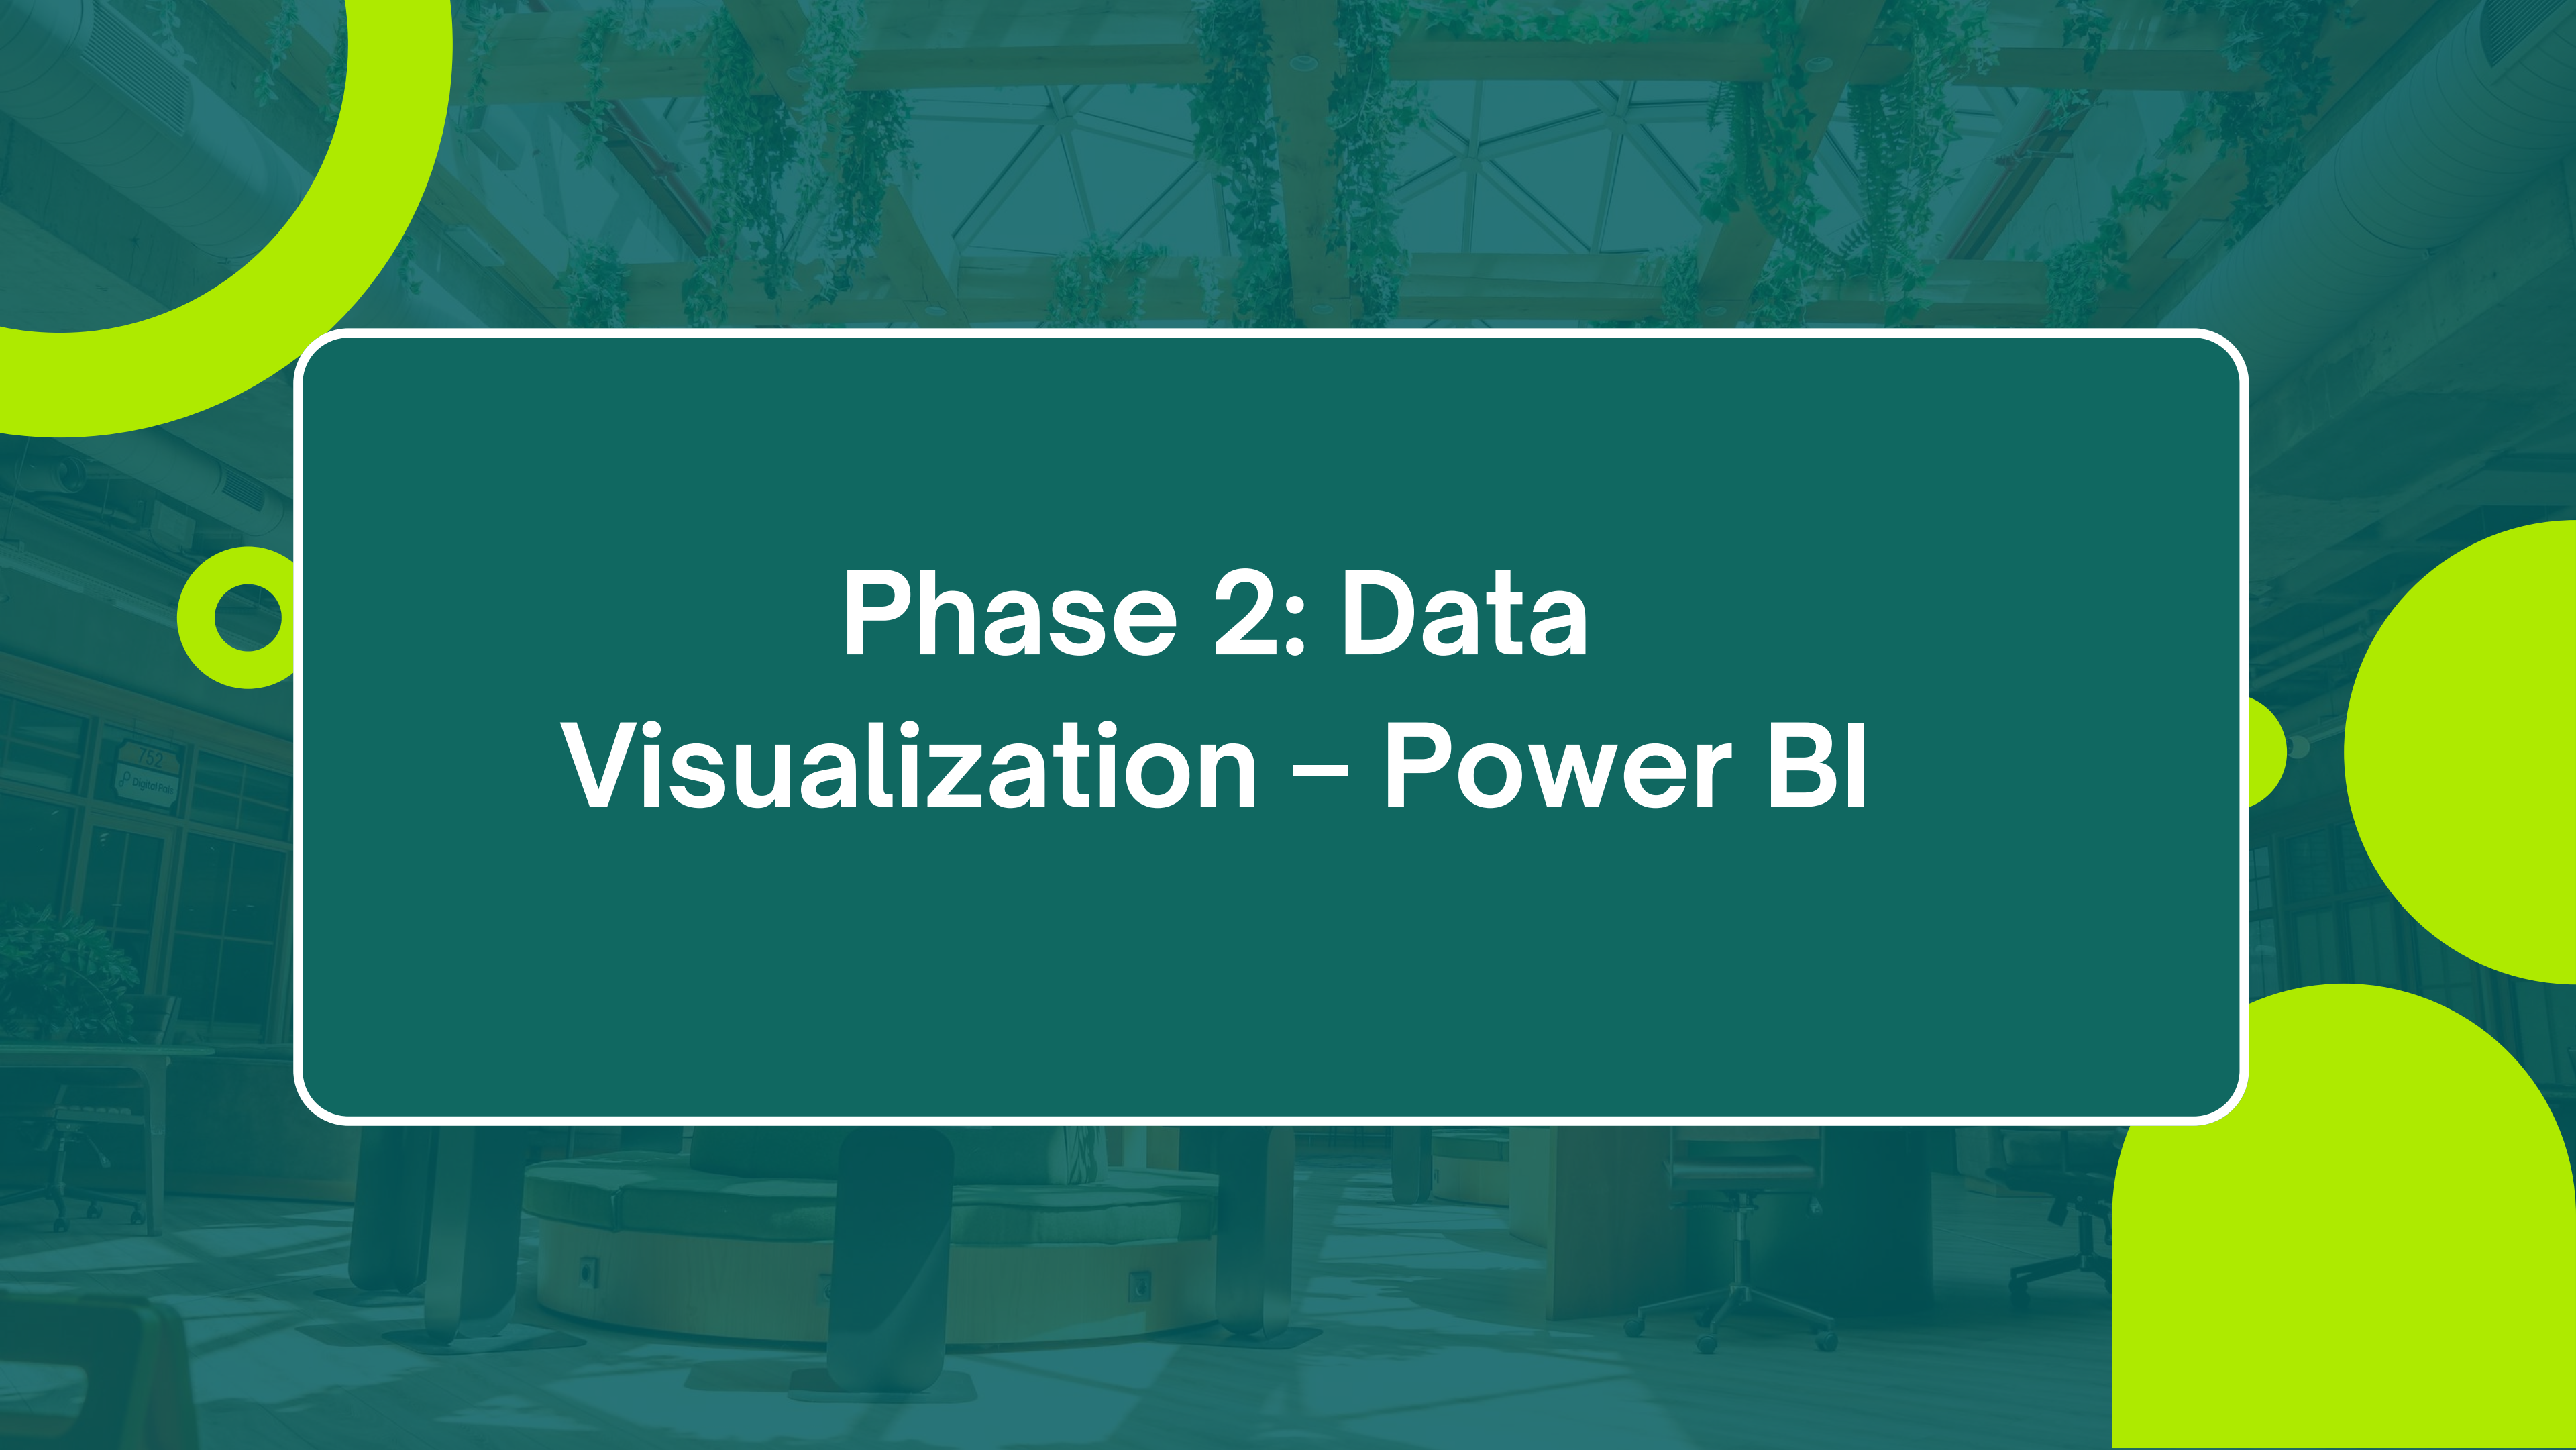
-- Combining MTD and YTD Metrics

```
SELECT
  YEAR(o.order_purchase_timestamp) AS year,
  DATE_FORMAT(o.order_purchase_timestamp, '%Y-%m') AS month,
  DATE(o.order_purchase_timestamp) AS order_date,
  SUM(p.payment_value) AS daily_sales,
  SUM(SUM(p.payment_value)) OVER (
    PARTITION BY YEAR(o.order_purchase_timestamp)
    ORDER BY DATE(o.order_purchase_timestamp)
  ) AS ytd_sales,
  SUM(SUM(p.payment_value)) OVER (
    PARTITION BY DATE_FORMAT(o.order_purchase_timestamp, '%Y-%m')
    ORDER BY DATE(o.order_purchase_timestamp)
  ) AS mtd_sales
FROM
  orders o
JOIN
  payments p
  ON o.order_id = p.order_id
WHERE
  o.order_status = 'delivered'
GROUP BY
  year, month, order_date
ORDER BY
  year, month, order_date;
```



| | year | month | order_date | daily_sales | ytd_sales | mtd_sales |
|---|------|---------|------------|--------------------|--------------------|--------------------|
| ▶ | 2016 | 2016-10 | 2016-10-03 | 70500.78128814697 | 70500.78128814697 | 70500.78128814697 |
| | 2016 | 2016-10 | 2016-10-04 | 1237498.9030365944 | 1307999.6843247414 | 1307999.6843247414 |
| | 2016 | 2016-10 | 2016-10-05 | 908396.9927902222 | 2216396.6771149635 | 2216396.6771149635 |
| | 2016 | 2016-10 | 2016-10-06 | 848423.5258255005 | 3064820.202940464 | 3064820.202940464 |
| | 2016 | 2016-10 | 2016-10-07 | 860927.7689208984 | 3925747.9718613625 | 3925747.9718613625 |
| | 2016 | 2016-10 | 2016-10-08 | 1067967.1884498596 | 4993715.160311222 | 4993715.160311222 |
| | 2016 | 2016-10 | 2016-10-09 | 382702.32545900345 | 5376417.4857702255 | 5376417.4857702255 |
| | 2016 | 2016-10 | 2016-10-10 | 490987.9774489403 | 5867405.463219166 | 5867405.463219166 |
| | 2016 | 2016-12 | 2016-12-23 | 2472.120105743408 | 5869877.583324909 | 2472.120105743408 |
| | 2017 | 2017-01 | 2017-01-05 | 89116.02258110046 | 89116.02258110046 | 89116.02258110046 |

Result 1 x



Phase 2: Data Visualization – Power BI

About Power BI

- **DAX Queries:** Used for creating complex calculations and data manipulations.
- **Visuals:** Allows users to create a wide range of charts and visuals (e.g., bar charts, line graphs, maps) to represent data insights.
- **Filters:** Enable users to segment data and focus on specific insights for better analysis.
- **Data Modeling:** Helps structure data relationships to ensure consistency and accuracy in reports.
- **Calendar Table:** Used for time-based analysis, supporting time intelligence functions like year-over-year comparisons and trend analysis.

Sales Performance Analysis:



\$13.22M

Total_Revenue

99K

Total Orders

\$133

Average Order Value

Average Order Value(AOV) by Month and Year

Year ● 2016 ● 2017 ● 2018



order_purchase_timestamp

All

product category

All


customer_state

All

Month

All


Total_Revenue by product category



Total_Revenue by product category

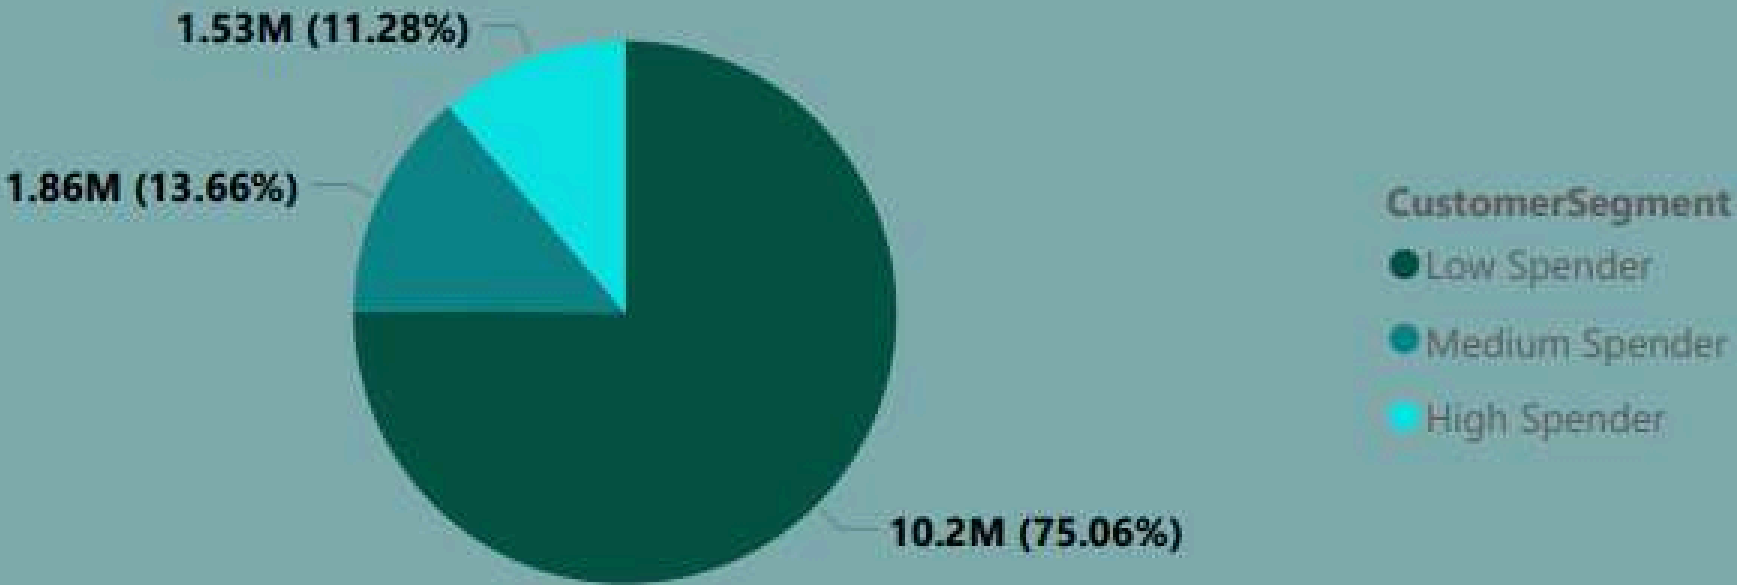


Customer Insights:



13.59M
CLV

TotalSpend by CustomerSegment



customer_state

All


order_purchase_timestamp

All

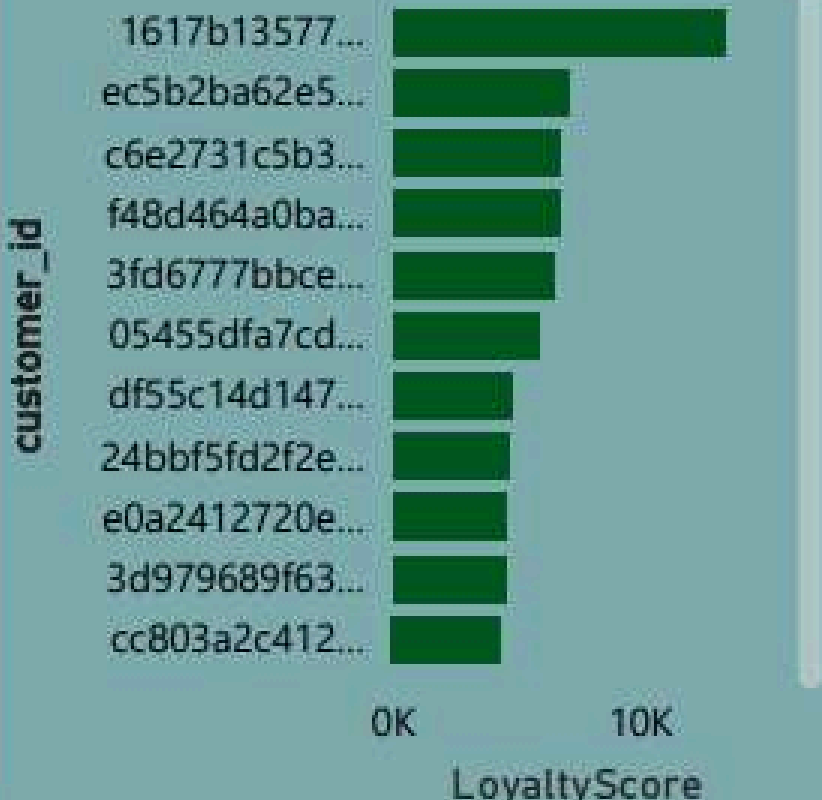
product category

All

TOP 10 customer by Revenue




LoyaltyScore by customer_id



| customer_id | CLV | CustomerSegment | PurchaseFrequency |
|----------------------------------|----------------|-----------------|-------------------|
| ffffe8b65bbe3087b653a978c870db99 | | No Purchase | 1 |
| ffffa3172527f765de70084a7e53aae8 | 21.80 | Low Spender | 1 |
| ffff42319e9b2d713724ae527742af25 | 199.90 | Low Spender | 1 |
| fffed5b6d849fbd39689bb92087f431 | 47.90 | Low Spender | 1 |
| fffecc9f79fd8c764f843e9951b11341 | 54.90 | Low Spender | 1 |
| fffc937e9dd47a13f05ecb8290f4d3e | 78.00 | Low Spender | 1 |
| fffc22669ca576ae3f654ea64c8f36be | | No Purchase | 1 |
| fffb97495f78be80e2759335275df2aa | 45.90 | Low Spender | 1 |
| fffa0238b217e18a8adeeda0669923a3 | 35.00 | Low Spender | 1 |
| fff93c1da78dafaaa304ff032abc6205 | 198.89 | Low Spender | 1 |
| Total | 1,35,91,643.70 | | 99441 |

Regional Analysis:



12.50

AvgDeliveryTimeByRegion

542

CanceledOrdersByProduct

0.63%

ReturnRateByRegion

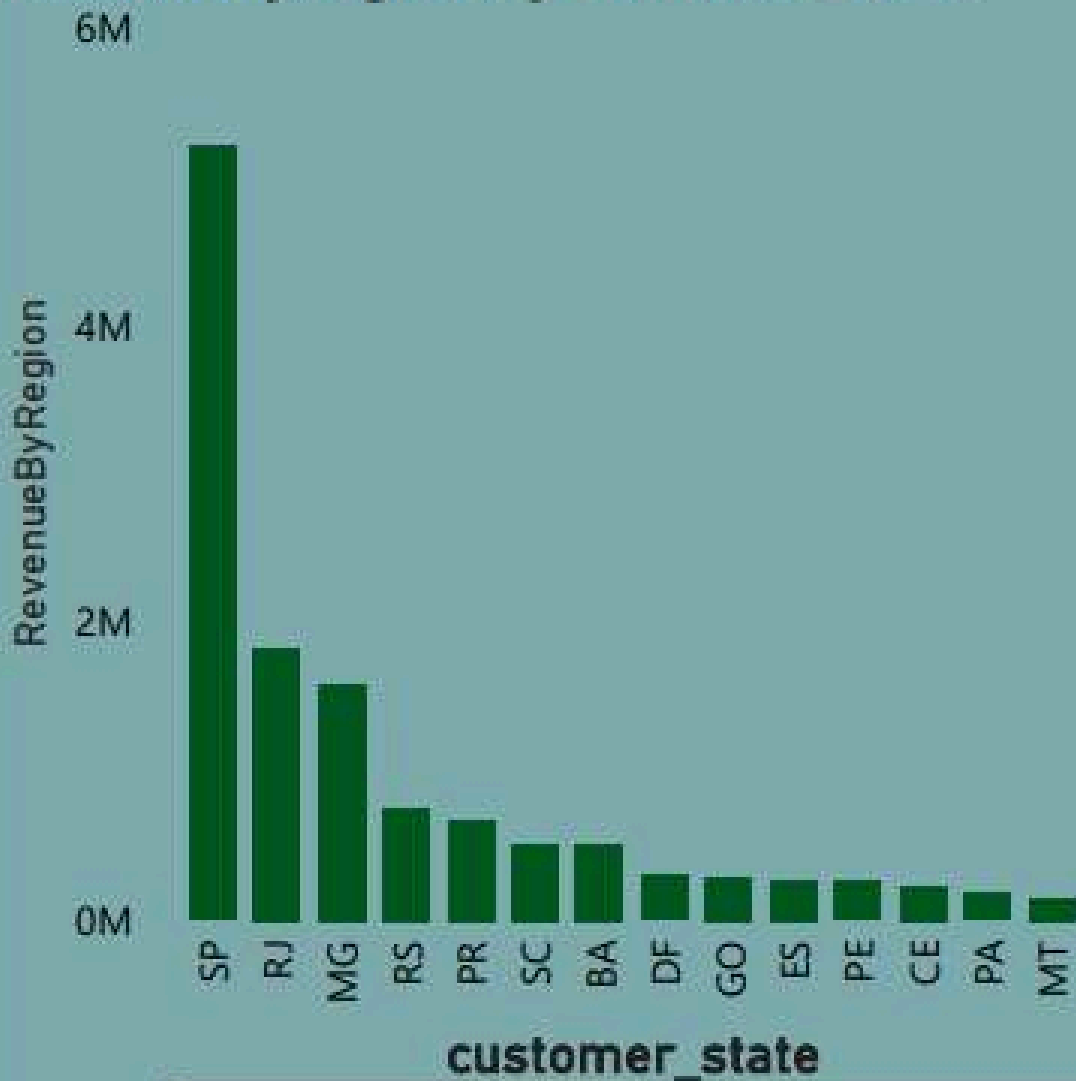
product category

All

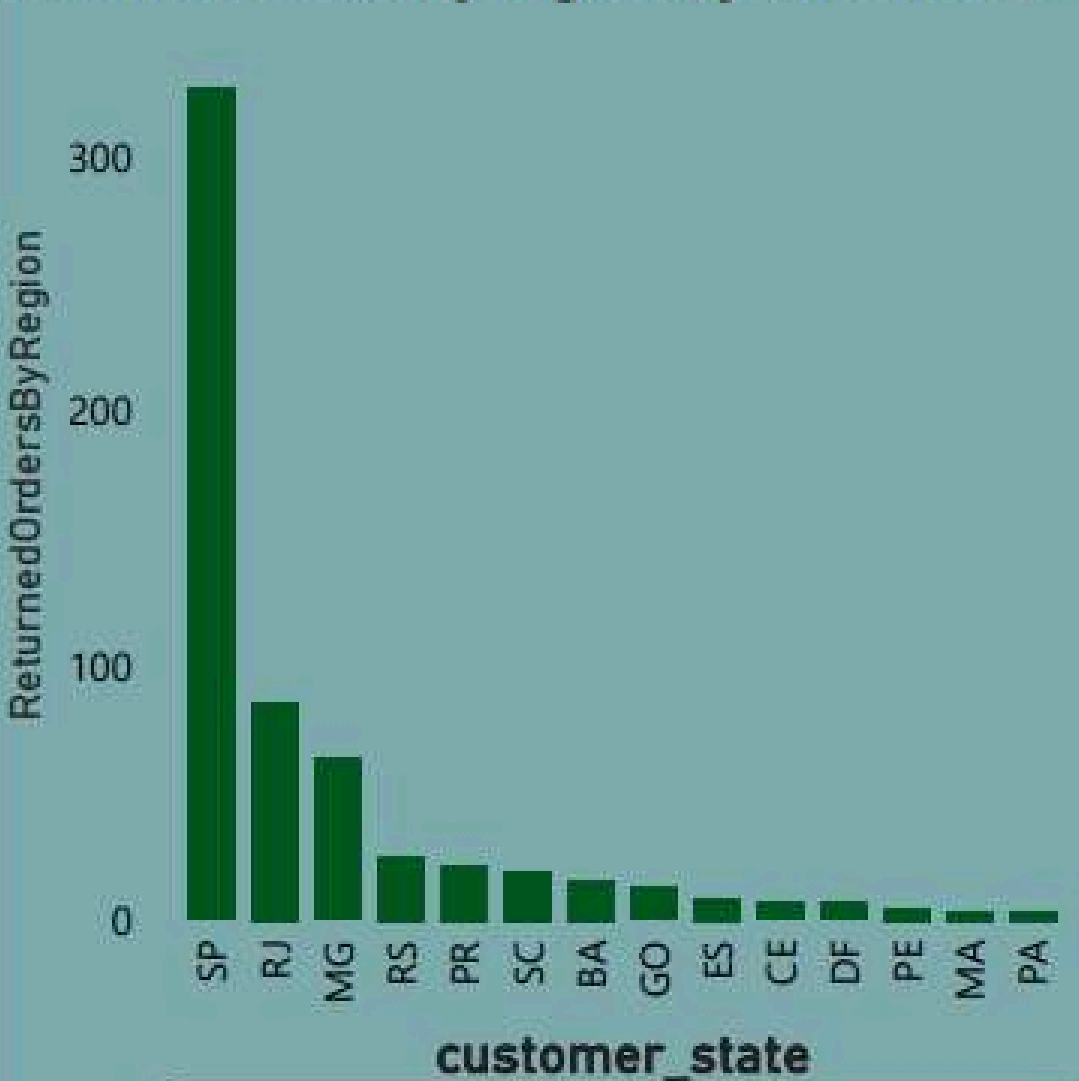
order_purchase_timestamp

All


RevenueByRegion by customer_state




ReturnedOrdersByRegion by customer_...



AvgDeliveryTimeByRegion by customer_state



Operational Metrics:



12.50

AvgDeliveryTime


210

MaxDeliveryTime

0

MinDeliveryTime

DeliveredOrdersByState by customer_state



order_purchase_timestamp

All

product category

All

customer_state

All


Month

All

| product category | TotalOrdersByCategory | RevenueByCategory |
|-------------------------|-----------------------|-------------------|
| Watches present | 5991 | 12,05,005.68 |
| toys | 4117 | 4,83,946.60 |
| telephony | 4545 | 3,23,667.53 |
| technical books | 267 | 19,096.06 |
| stationary store | 2517 | 2,30,943.23 |
| sport leisure | 8641 | 9,88,048.97 |
| song | 38 | 6,034.35 |
| SIGNIFICATION AND CASES | 188 | 31,588.23 |
| Total | 112650 | 1,35,91,643.70 |

| product category | CanceledOrdersByProduct |
|------------------|-------------------------|
| Watches present | 21 |
| toys | 34 |
| telephony | 18 |
| stationary store | 12 |
| sport leisure | 51 |
| Room Furniture | 2 |
| Total | 542 |

Time Trends:



\$13.59M

TotalRevenue

0%

MoM_Growth

250.3%

YoY_Growth

TotalRevenue by Month and Year

Year: ● 2016 ● 2017 ● 2018



Year

All

Month

All

customer_state

All

| Year | TotalRevenue | YoY_Growth |
|-------|----------------|------------|
| 2017 | 55,17,496.55 | 132.98 |
| 2018 | 80,24,098.21 | 1.10 |
| | 893.90 | 0.00 |
| 2016 | 49,155.04 | 0.00 |
| Total | 1,35,91,643.70 | 2.50 |

| product category | MoM_Growth | YoY_Growth | TotalRevenue |
|------------------|------------|------------|----------------|
| Watches present | 0.00 | 3.42 | 12,05,005.68 |
| toys | 0.00 | 1.86 | 4,83,946.60 |
| telephony | 0.00 | 2.89 | 3,23,667.53 |
| technical books | 0.01 | 4.38 | 19,096.06 |
| stationary store | 0.00 | 4.15 | 2,30,943.23 |
| sport leisure | 0.00 | 2.43 | 9,88,048.97 |
| | 0.00 | 2.01 | 6,024.25 |
| Total | 0.00 | 2.50 | 1,35,91,643.70 |

Key Findings

Revenue Insights:

Total revenue: \$13.22M, with an Average Order Value (AOV) of \$133.04.

Health and Beauty is the top product category, contributing the most to revenue.

Customer Insights:

Loyal customers show high purchase frequency and contribute significantly to revenue.

The top 10 customers were identified using the Loyalty Factor, combining purchase frequency and spending metrics.

Operational Efficiency:

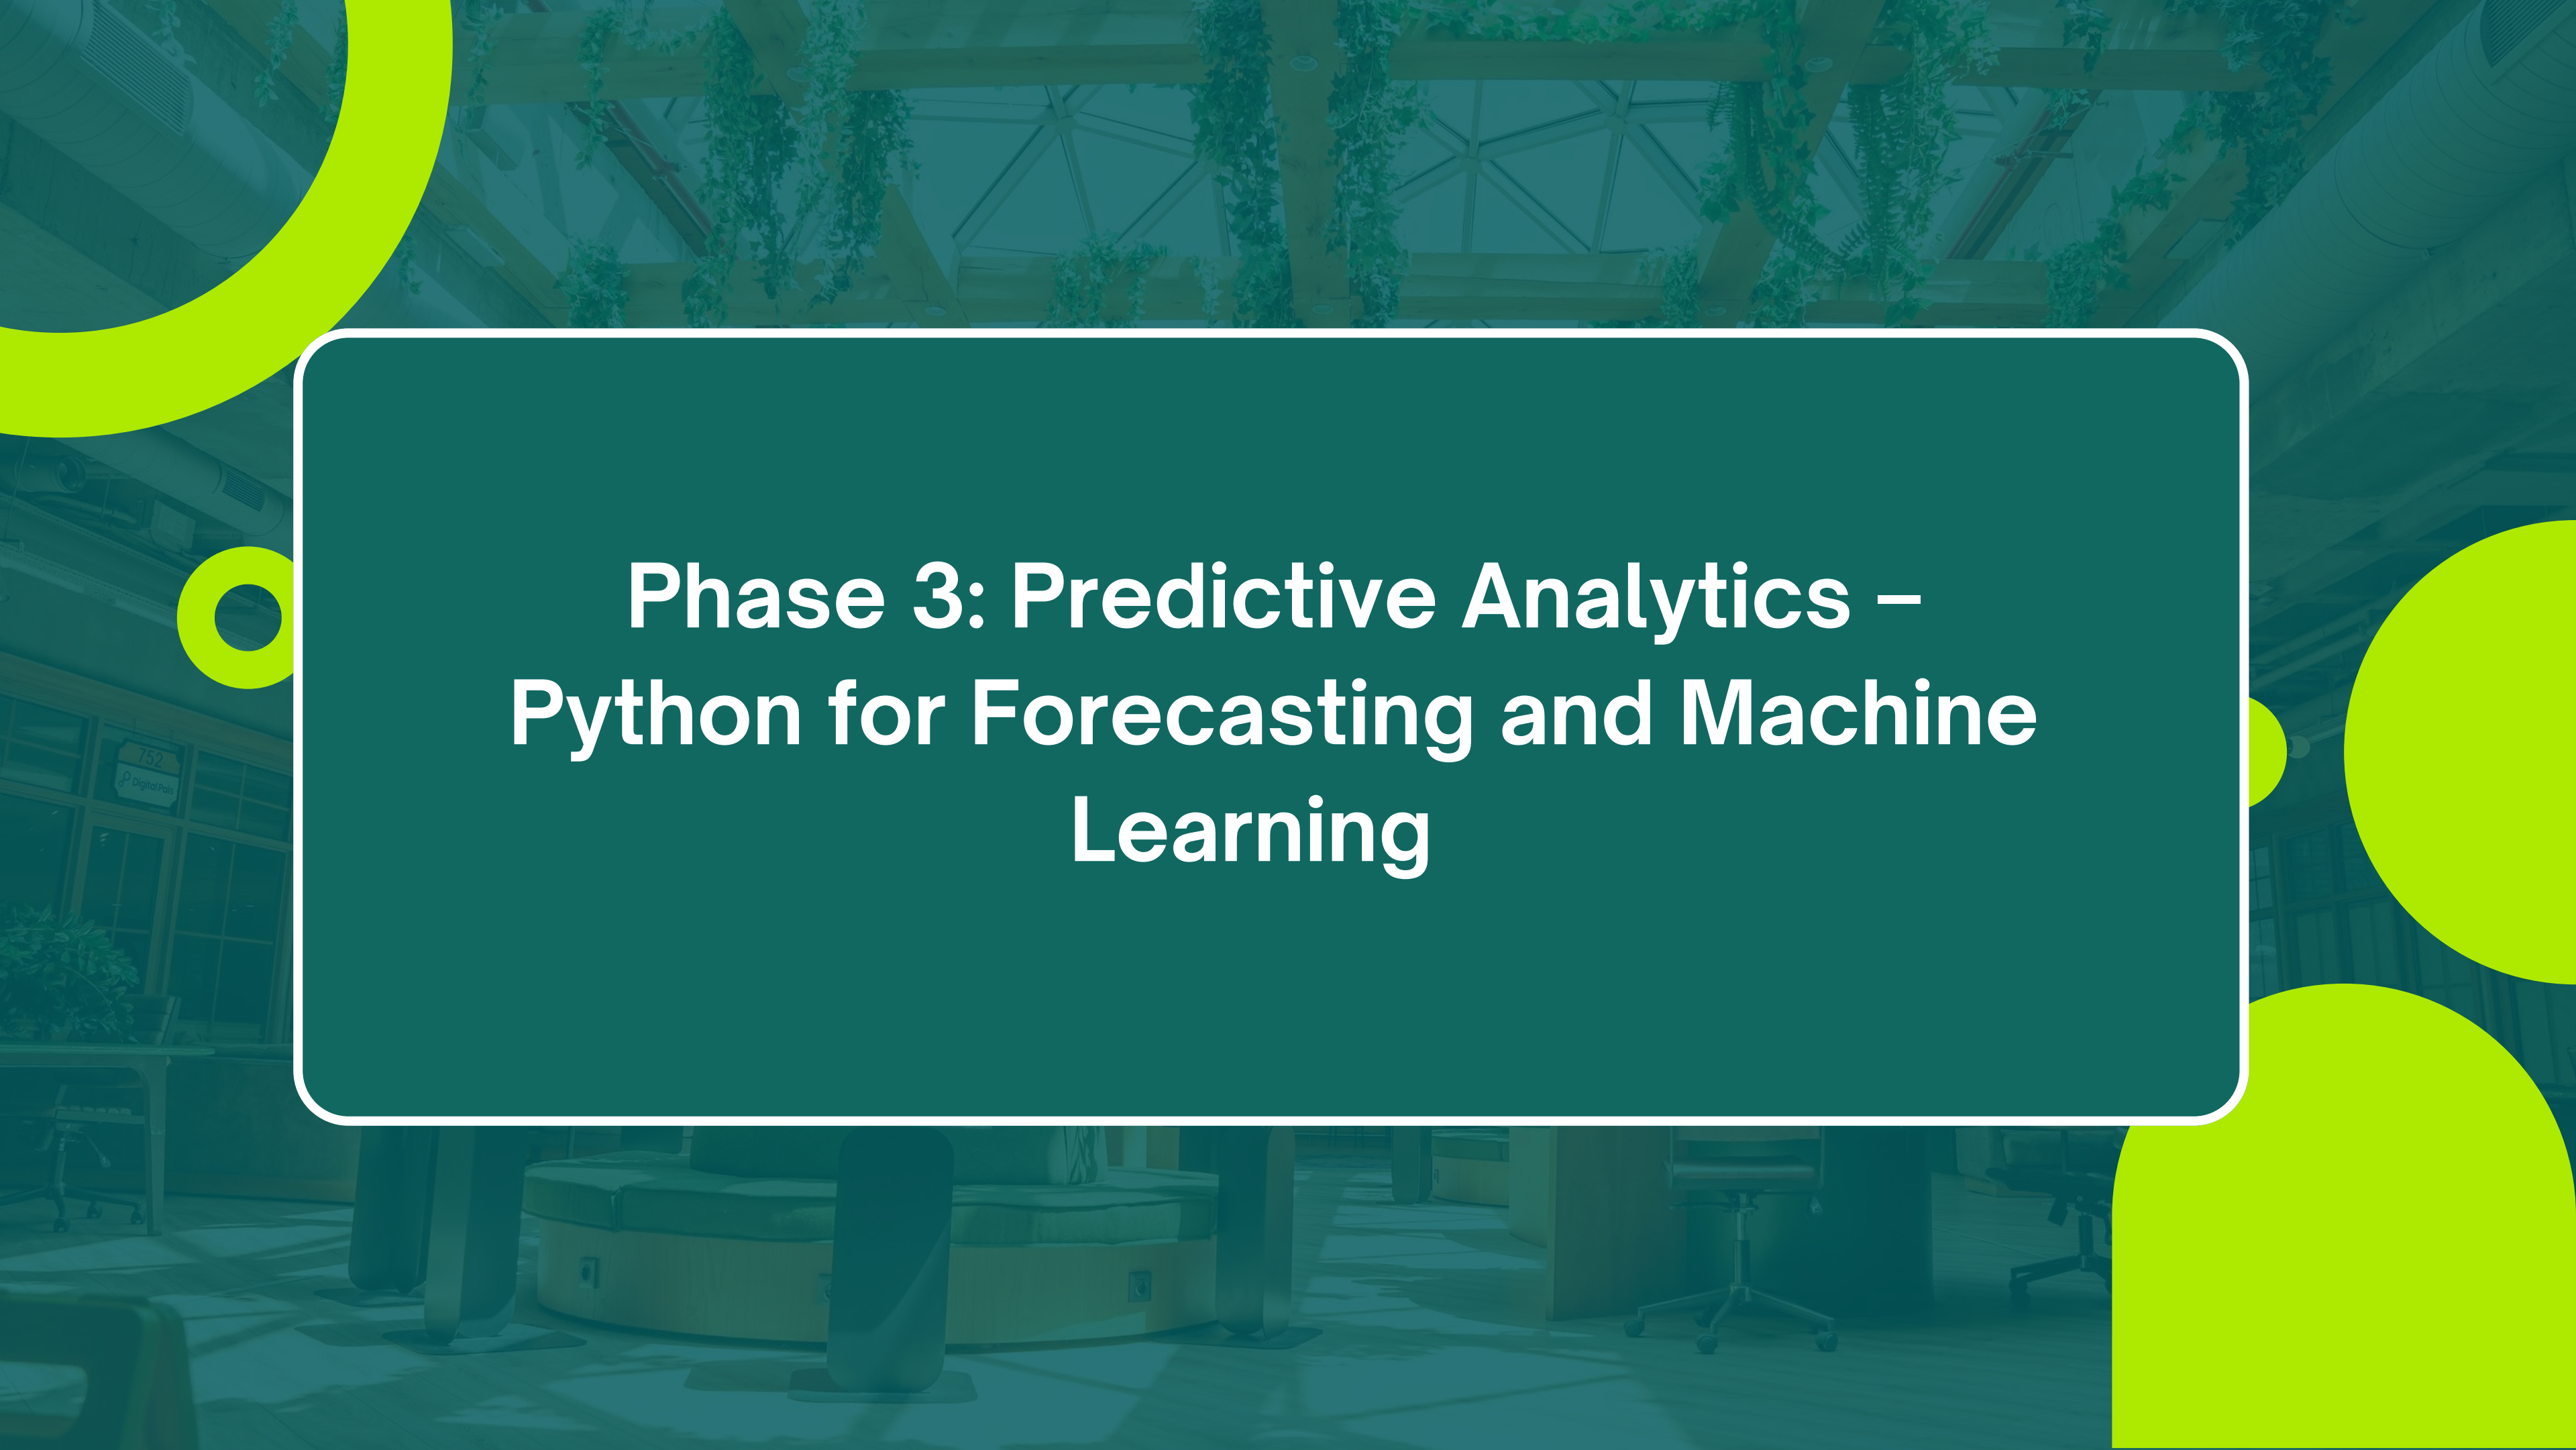
Average delivery time: 12.50 days, with delays in some regions.

Peak sales months and weekday/weekend patterns identified for targeted planning.

Order Insights:

Total products ordered: 99K, with 542 cancellations, highlighting areas for improvement.

Strategic seasonal marketing campaigns during peak months to maximize sales.



Phase 3: Predictive Analytics – Python for Forecasting and Machine Learning

About Python

Python Libraries

1. Pandas

- **Purpose:** Data manipulation and analysis.
- **Usage:** Used for reading, merging, filtering, and handling datasets (e.g., merged_items DataFrame).

2. NumPy

- **Purpose:** Numerical computing.
- **Usage:** For creating arrays, handling numerical data, and initializing binary data for the sparse matrix.

3. SciPy

- **Module Used:** scipy.sparse.coo_matrix
- **Purpose:** Efficient storage and computation with sparse matrices.
- **Usage:** Used to construct a binary sparse matrix to represent the product-order relationship.

4. scikit-learn

- **Module Used:** sklearn.metrics.pairwise.cosine_similarity
- **Purpose:** Machine learning and statistical tools.
- **Usage:** Used for computing pairwise cosine similarity to identify product correlations.

Development Environment

• Jupyter Notebook

- **Purpose:** Interactive environment for writing, testing, and visualizing Python code.
- **Usage:** Used as the primary platform to execute and document the code.

Tasks and Techniques:

Q1. Sales Prediction:

- Use linear regression to predict monthly sales for the next 12 months.
- Identify key factors (e.g., past sales, promotions, and external variables) influencing sales trends

Step 1: Define the Problem

- We aim to predict monthly sales for the next 12 months using linear regression. We will:
- Use the orders.csv and order_items.csv tables to calculate historical sales.
- Aggregate sales data to derive monthly sales.
- Use the historical data as a basis to train a linear regression model.

Step 2: Prepare the Data

a) Load the necessary tables

- We will focus on the following tables:
- orders.csv: To extract the order dates (order_purchase_timestamp) and associate them with order IDs. order_items.csv: To get the sales amount for each order using the price column.

```
import pandas as pd
```

Load datasets

```
orders = pd.read_csv("E:\BHU\Data_Analyst\Ecommerce\orders.csv")  
order_items = pd.read_csv("E:\BHU\Data_Analyst\Ecommerce\order_items.csv")
```

Display initial datasets

```
print(orders.head())  
print(order_items.head())
```

b) Merge and clean the data

- Join orders and order_items on the order_id column to combine order dates with sales data.

Merge orders with order_items

```
merged_data = pd.merge(order_items, orders, on="order_id")
```

Ensure timestamps are in datetime format

```
merged_data['order_purchase_timestamp'] = pd.to_datetime(merged_data['order_purchase_timestamp'])
```

Drop rows with missing order dates or sales amounts

```
merged_data = merged_data.dropna(subset=['order_purchase_timestamp', 'price'])
```

Display the cleaned merged dataset

```
print(merged_data.head())
```

c) Aggregate monthly sales

- Aggregate the sales data (price) by month and calculate total monthly sales.

Extract month and year

```
merged_data['month_year'] = merged_data['order_purchase_timestamp'].dt.to_period('M')
```

Calculate monthly sales

```
monthly_sales = merged_data.groupby('month_year')['price'].sum().reset_index()  
monthly_sales.rename(columns={'price': 'total_sales'}, inplace=True)
```

Convert month_year back to datetime for plotting and modeling

```
monthly_sales['month_year'] = monthly_sales['month_year'].astype('datetime64[ns]')
```

Display monthly sales data

```
print(monthly_sales.head())
```

Output:-

```
   month_year  total_sales  
0 2016-09-01      267.36  
1 2016-10-01    49507.66  
2 2016-12-01       10.90  
3 2017-01-01   120312.87  
4 2017-02-01   247303.02
```


Step 3: Visualize the Data

- Visualize historical monthly sales to understand trends and seasonality.

```
import matplotlib.pyplot as plt
```

```
# Plot the historical monthly sales
```

```
plt.figure(figsize=(12, 6))
```

```
plt.plot(monthly_sales['month_year'], monthly_sales['total_sales'], marker='o', label="Historical Sales")
```

```
plt.title("Monthly Sales Over Time")
```

```
plt.xlabel("Month")
```

```
plt.ylabel("Total Sales")
```

```
plt.grid(True)
```

```
plt.legend()
```

```
plt.show()
```

Output:-



Step 4: Build the Linear Regression Model

a) Prepare the features and target variable

Create a time index (X) and corresponding sales (y).

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import numpy as np
```

Create time index as a numerical feature

```
monthly_sales['time_index'] = np.arange(len(monthly_sales))
```

Features (time index) and target (sales)

```
X = monthly_sales[['time_index']] # Predictor
```

```
y = monthly_sales['total_sales'] # Target
```

Split the data into training and test sets (80%-20%)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

b) Train the model

Train a linear regression model using the training data.

Initialize and train the model

```
model = LinearRegression()  
model.fit(X_train, y_train)
```

Display the model coefficients

```
print(f"Intercept: {model.intercept_}")  
print(f"Slope: {model.coef_[0]}")
```

Output:-

```
Intercept: 188578.81335401285  
Slope: 31883.910924994423
```

b) Evaluate the model

Check the model's performance on the test set.

```
from sklearn.metrics import mean_squared_error, r2_score
```

Predict on the test set

```
y_pred = model.predict(X_test)
```

Calculate evaluation metrics

```
mse = mean_squared_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)
```

```
print(f"Mean Squared Error: {mse}")  
print(f"R^2 Score: {r2}")
```

Output:-

```
Mean Squared Error: 23760769463.13549  
R^2 Score: 0.8029932021633788
```


Step 5: Forecast Sales for the Next 12 Months

a) Extend the time index

Generate future time indices for the next 12 months.

Get the last time index

```
last_time_index = monthly_sales['time_index'].max()
```

Create future time indices

```
future_time_indices = np.arange(last_time_index + 1, last_time_index + 13).reshape(-1, 1)
```

b) Predict future sales

Use the trained model to forecast sales.

Predict future sales

```
future_sales = model.predict(future_time_indices)
```

Create a DataFrame for future predictions

```
future_months = pd.date_range(start=monthly_sales['month_year'].max() + pd.offsets.MonthBegin(1),  
periods=12, freq='M')
```

```
future_sales_df = pd.DataFrame({  
    'month_year': future_months,  
    'predicted_sales': future_sales  
})
```

Display future sales predictions

```
print(future_sales_df)
```

Output:-

| | month_year | predicted_sales |
|----|------------|-----------------|
| 0 | 2018-10-31 | 9.537927e+05 |
| 1 | 2018-11-30 | 9.856766e+05 |
| 2 | 2018-12-31 | 1.017560e+06 |
| 3 | 2019-01-31 | 1.049444e+06 |
| 4 | 2019-02-28 | 1.081328e+06 |
| 5 | 2019-03-31 | 1.113212e+06 |
| 6 | 2019-04-30 | 1.145096e+06 |
| 7 | 2019-05-31 | 1.176980e+06 |
| 8 | 2019-06-30 | 1.208864e+06 |
| 9 | 2019-07-31 | 1.240748e+06 |
| 10 | 2019-08-31 | 1.272632e+06 |
| 11 | 2019-09-30 | 1.304516e+06 |

C:\Users\De11\anaconda3\lib\site-packages

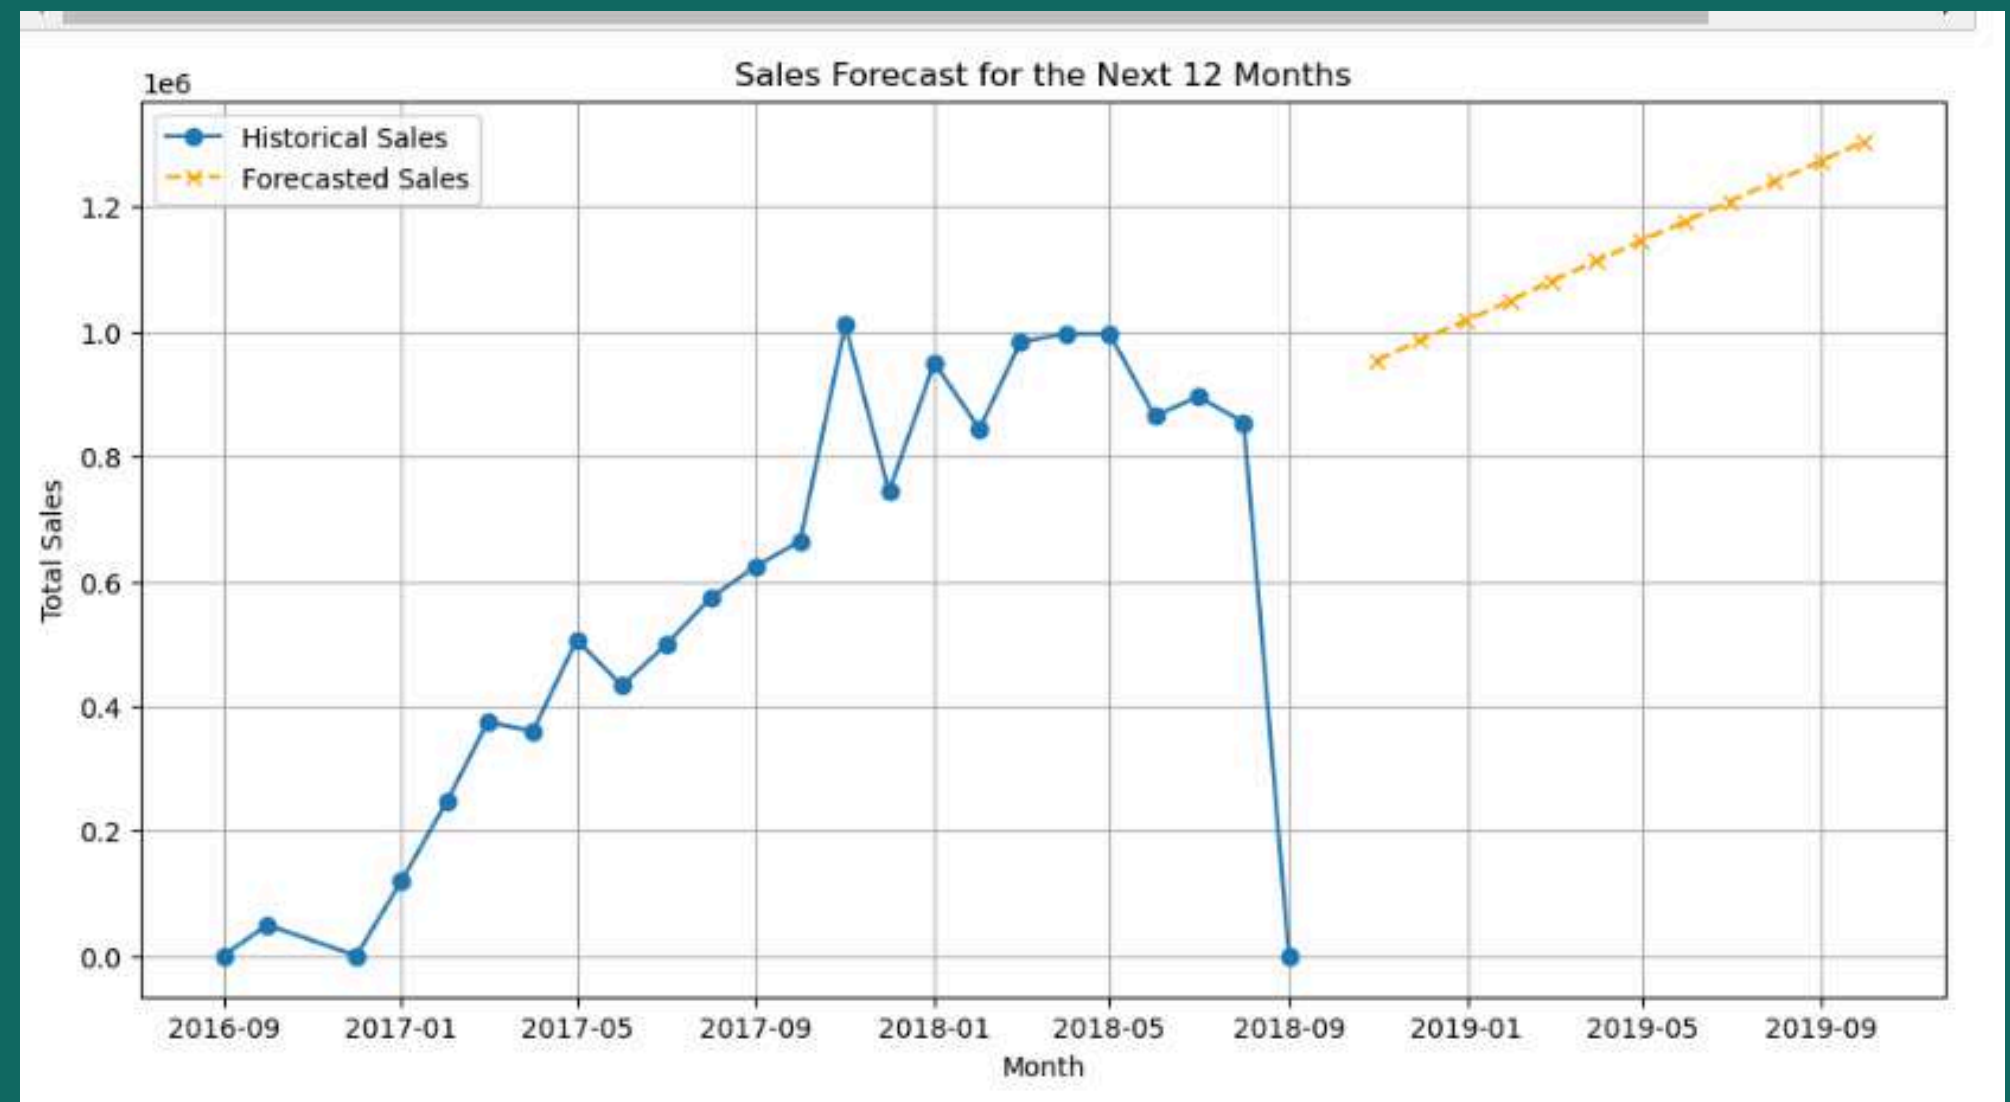
c) Visualize the forecast

Combine historical sales with forecasted sales for visualization.

Plot historical and forecasted sales

```
plt.figure(figsize=(12, 6))
plt.plot(monthly_sales['month_year'], monthly_sales['total_sales'], marker='o', label="Historical Sales")
plt.plot(future_sales_df['month_year'], future_sales_df['predicted_sales'], marker='x', linestyle='--',
color='orange', label="Forecasted Sales")
plt.title("Sales Forecast for the Next 12 Months")
plt.xlabel("Month")
plt.ylabel("Total Sales")
plt.grid(True)
plt.legend()
plt.show()
```

Output:-



Q2. Customer Segmentation (Clustering):

Perform RFM Analysis (Recency, Frequency, Monetary) and use simple clustering methods like hierarchical clustering or manual thresholds to classify customers.

Label customer segments (e.g., Loyal Customers, Discount Seekers) to identify marketing opportunities.

Step 1: Define the Objective

We aim to segment customers using RFM Analysis:

Recency: How recently a customer made a purchase.

Frequency: How often a customer made purchases.

Monetary: How much a customer spent.

Step 2: Prepare the Data

a) Load the necessary tables

Load the orders.csv and order_items.csv tables.

```
import pandas as pd
```

Load datasets

```
orders = pd.read_csv("E:\BHU\Data_Analyst\Ecommerce\orders.csv")
```

```
order_items = pd.read_csv("E:\BHU\Data_Analyst\Ecommerce\order_items.csv")
```

Display the first few rows of the datasets

```
print(orders.head())
```

```
print(order_items.head())
```

Output:-

```
order_id      customer_id \
0  e481f51cbdc54678b7cc49136f2d6af7  9ef432eb6251297304e76186b10a928d
1  53cdb2fc8bc7dce0b6741e2150273451  b0830fb4747a6c6d20dea0b8c802d7ef
2  47770eb9100c2d0c44946d9cf07ec65d  41ce2a54c0b03bf3443c3d931a367089
3  949d5b44dbf5de918fe9c16f97b45f8a  f88197465ea7920adcdbe7375364d82
4  ad21c59c0840e6cb83a9ceb5573f8159  8ab97904e6daea8866dbdbc4fb7aad2c

order_status  order_purchase_timestamp  order_approved_at \
0  delivered      2017-10-02 10:56:33      2017-10-02 11:07:15
1  delivered      2018-07-24 20:41:37      2018-07-26 03:24:27
2  delivered      2018-08-08 08:38:49      2018-08-08 08:55:23
3  delivered      2017-11-18 19:28:06      2017-11-18 19:45:59
4  delivered      2018-02-13 21:18:39      2018-02-13 22:20:29

order_delivered_carrier_date  order_delivered_customer_date \
0      2017-10-04 19:55:00      2017-10-10 21:25:13
1      2018-07-26 14:31:00      2018-08-07 15:27:45
2      2018-08-08 13:50:00      2018-08-17 18:06:29
3      2017-11-22 13:39:59      2017-12-02 00:28:42
4      2018-02-14 19:46:34      2018-02-16 18:17:02

order_estimated_delivery_date
0      2017-10-18 00:00:00
1      2018-08-13 00:00:00
2      2018-09-04 00:00:00
3      2017-12-15 00:00:00
4      2018-02-26 00:00:00

order_id      order_item_id \
0  00010242fe8c5a6d1ba2dd792cb16214      1
1  00018f77f2f0320c557190d7a144bdd3      1
2  000229ec398224ef6ca0657da4fc703e      1
3  00024acbcd0a6daa1e931b038114c75      1
4  00042b26cf59d7ce69dfabb4e55b4fd9      1

product_id      seller_id \
0  4244733e06e7ecb4970a6e2683c13e61  48436dade18ac8b2bce089ec2a041202
1  e5f2d52b002189ee658865ca93d83a8f  dd7ddc04e1b6c2c614352b383efe2d36
2  c777355d18b72b67abbef9df44fd0fd  5b51032eddd242adc84c38acab88f23d
3  7634da152a4610f1595efa32f14722fc  9d7a1d34a5052409006425275ba1c2b4
4  ac6c3623068f30de03045865e4e10089  df560393f3a51e74553ab94004ba5c87

shipping_limit_date  price  freight_value
0  2017-09-19 09:45:35    58.90          13.29
1  2017-05-03 11:05:13   239.90          19.93
2  2018-01-18 14:48:30   199.00          17.87
3  2018-08-15 10:10:18    12.99          12.79
4  2017-02-13 13:57:51   199.90          18.14
```


b) Merge and clean the data

Join orders and order_items on the order_id column to combine customer information with sales data.

Merge orders with order_items

```
rfm_data = pd.merge(order_items, orders, on="order_id")
```

Ensure timestamps are in datetime format

```
rfm_data['order_purchase_timestamp'] = pd.to_datetime(rfm_data['order_purchase_timestamp'])
```

Drop rows with missing values

```
rfm_data = rfm_data.dropna(subset=['order_purchase_timestamp', 'price', 'customer_id'])
```

Display the merged dataset

```
print(rfm_data.head())
```

Output:-

| | order_id | order_item_id | \ |
|---|----------------------------------|---------------|---|
| 0 | 00010242fe8c5a6d1ba2dd792cb16214 | 1 | |
| 1 | 00018f77f2f0320c557190d7a144bdd3 | 1 | |
| 2 | 000229ec398224ef6ca0657da4fc703e | 1 | |
| 3 | 00024acbcd0a6daa1e931b038114c75 | 1 | |
| 4 | 00042b26cf59d7ce69dfabb4e55b4fd9 | 1 | |

| | product_id | seller_id | \ |
|---|----------------------------------|----------------------------------|---|
| 0 | 4244733e06e7ecb4970a6e2683c13e61 | 48436dade18ac8b2bce089ec2a041202 | |
| 1 | e5f2d52b802189ee658865ca93d83a8f | dd7ddc04e1b6c2c614352b383efe2d36 | |
| 2 | c777355d18b72b67abbef9df44fd0fd | 5b51032eddd242adc84c38acab88f23d | |
| 3 | 7634da152a4610f1595efa32f14722fc | 9d7a1d34a5052409006425275ba1c2b4 | |
| 4 | ac6c3623068f30de03045865e4e10089 | df560393f3a51e74553ab94004ba5c87 | |

| | shipping_limit_date | price | freight_value | \ |
|---|---------------------|--------|---------------|---|
| 0 | 2017-09-19 09:45:35 | 58.90 | 13.29 | |
| 1 | 2017-05-03 11:05:13 | 239.90 | 19.93 | |
| 2 | 2018-01-18 14:48:30 | 199.00 | 17.87 | |
| 3 | 2018-08-15 10:10:18 | 12.99 | 12.79 | |
| 4 | 2017-02-13 13:57:51 | 199.90 | 18.14 | |

| | customer_id | order_status | order_purchase_timestamp | \ |
|---|----------------------------------|--------------|--------------------------|---|
| 0 | 3ce436f183e68e07877b285a838db11a | delivered | 2017-09-13 08:59:02 | |
| 1 | f6dd3ec061db4e3987629fe6b26e5cce | delivered | 2017-04-26 10:53:06 | |
| 2 | 6489ae5e4333f3693df5ad4372dab6d3 | delivered | 2018-01-14 14:33:31 | |
| 3 | d4eb9395c8c0431ee92fce09860c5a06 | delivered | 2018-08-08 10:00:35 | |
| 4 | 58dbd0b2d70206bf40e62cd34e84d795 | delivered | 2017-02-04 13:57:51 | |

| | order_approved_at | order_delivered_carrier_date | \ |
|---|---------------------|------------------------------|---|
| 0 | 2017-09-13 09:45:35 | 2017-09-19 18:34:16 | |
| 1 | 2017-04-26 11:05:13 | 2017-05-04 14:35:00 | |
| 2 | 2018-01-14 14:48:30 | 2018-01-16 12:36:48 | |
| 3 | 2018-08-08 10:10:18 | 2018-08-10 13:28:00 | |
| 4 | 2017-02-04 14:10:13 | 2017-02-16 09:46:09 | |

| | order_delivered_customer_date | order_estimated_delivery_date |
|---|-------------------------------|-------------------------------|
| 0 | 2017-09-20 23:43:48 | 2017-09-29 00:00:00 |
| 1 | 2017-05-12 16:04:24 | 2017-05-15 00:00:00 |
| 2 | 2018-01-22 13:19:16 | 2018-02-05 00:00:00 |
| 3 | 2018-08-14 13:32:39 | 2018-08-20 00:00:00 |
| 4 | 2017-03-01 16:42:31 | 2017-03-17 00:00:00 |

Step 3: Compute RFM Metrics

a) Define the reference date Choose a reference date (e.g., the most recent date in the dataset).

Reference date for recency calculation

```
reference_date = rfm_data['order_purchase_timestamp'].max()
print(f"Reference Date: {reference_date}")
```

Output:-

```
Reference Date: 2018-09-03 09:06:57
```

b) Calculate RFM metrics

Group the data by customer_id to calculate Recency, Frequency, and Monetary values.

```
# Group by customer_id
rfm_metrics = rfm_data.groupby('customer_id').agg({
    'order_purchase_timestamp': lambda x: (reference_date - x.max()).days, # Recency
    'order_id': 'count', # Frequency
    'price': 'sum' # Monetary
}).reset_index()
```

Rename the columns

```
rfm_metrics.rename(columns={
    'order_purchase_timestamp': 'Recency',
    'order_id': 'Frequency',
    'price': 'Monetary'
}, inplace=True)
```

Display the RFM metrics

```
print(rfm_metrics.head())
```

Output:-

| | customer_id | Recency | Frequency | Monetary |
|---|-----------------------------------|---------|-----------|----------|
| 0 | 00012a2ce6f8dcda20d059ce98491703 | 292 | 1 | 89.80 |
| 1 | 000161a058600d5901f007fab4c27140 | 413 | 1 | 54.90 |
| 2 | 0001fd6190edaaaf884bcaf3d49edf079 | 551 | 1 | 179.99 |
| 3 | 0002414f95344307404f0ace7a26f1d5 | 382 | 1 | 149.90 |
| 4 | 000379cdec625522490c315e70c7a9fb | 153 | 1 | 93.00 |

Step 4: Perform Clustering

a) Normalize the data

Since RFM values have different scales, normalize them for clustering.

```
from sklearn.preprocessing import MinMaxScaler
```

```
# Initialize a scaler
```

```
scaler = MinMaxScaler()
```

```
# Normalize the RFM metrics
```

```
rfm_normalized = scaler.fit_transform(rfm_metrics[['Recency', 'Frequency', 'Monetary']])
```

```
# Convert back to DataFrame
```

```
rfm_normalized = pd.DataFrame(rfm_normalized, columns=['Recency', 'Frequency', 'Monetary'])
```

```
print(rfm_normalized.head())
```

Output:-

| | Recency | Frequency | Monetary |
|---|----------|-----------|----------|
| 0 | 0.401099 | 0.0 | 0.006619 |
| 1 | 0.567308 | 0.0 | 0.004022 |
| 2 | 0.756868 | 0.0 | 0.013330 |
| 3 | 0.524725 | 0.0 | 0.011091 |
| 4 | 0.210165 | 0.0 | 0.006857 |

b) Switch to K-Means Clustering

K-Means clustering is computationally efficient and can handle large datasets without requiring a pairwise distance matrix.

```
from sklearn.cluster import KMeans
import numpy as np
```

```
# Choose the number of clusters (e.g., 4 clusters)
```

```
n_clusters = 4
```

```
# Perform K-Means clustering
```

```
kmeans = KMeans(n_clusters=n_clusters, random_state=42)
```

```
rfm_metrics['Cluster'] = kmeans.fit_predict(rfm_normalized)
```

```
# Display the cluster assignments
```

```
print(rfm_metrics[['customer_id', 'Recency', 'Frequency', 'Monetary', 'Cluster']].head())
```

Output:-

```
C:\Users\Dell\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: `warn` will be removed in 1.4. Set the value of `n_init` explicitly to suppress warnings.warn(
```

| | customer_id | Recency | Frequency | Monetary | Cluster |
|---|-----------------------------------|---------|-----------|----------|---------|
| 0 | 00012a2ce6f8dcda20d059ce98491703 | 292 | 1 | 89.80 | 3 |
| 1 | 000161a058600d5901f007fab4c27140 | 413 | 1 | 54.90 | 1 |
| 2 | 0001fd6190edaaf884bcacf3d49edf079 | 551 | 1 | 179.99 | 1 |
| 3 | 0002414f95344307404f0ace7a26f1d5 | 382 | 1 | 149.90 | 3 |
| 4 | 000379cdec625522490c315e70c7a9fb | 153 | 1 | 93.00 | 2 |

Step 5: Label Customer Segments

Label the clusters based on RFM values (e.g., Loyal Customers, Discount Seekers).

Define segments based on cluster analysis

```
def label_segment(row):
    if row['Cluster'] == 0:
        return "Loyal Customers"
    elif row['Cluster'] == 1:
        return "Big Spenders"
    elif row['Cluster'] == 2:
        return "Occasional Buyers"
    elif row['Cluster'] == 3:
        return "At-Risk Customers"
    else:
        return "Other"
```

Apply labels to the clusters

```
rfm_metrics['Segment'] = rfm_metrics.apply(label_segment, axis=1)
```

Display customer segments

```
print(rfm_metrics[['customer_id', 'Recency', 'Frequency', 'Monetary', 'Cluster', 'Segment']].head())
```

Output:-

| | customer_id | Recency | Frequency | Monetary | Cluster | \ |
|---|----------------------------------|---------|-----------|----------|---------|---|
| 0 | 00012a2ce6f8dcda20d059ce98491703 | 292 | 1 | 89.80 | 3 | |
| 1 | 000161a058600d5901f007fab4c27140 | 413 | 1 | 54.90 | 1 | |
| 2 | 0001fd6190edaaf884bcaf3d49edf079 | 551 | 1 | 179.99 | 1 | |
| 3 | 0002414f95344307404f0ace7a26f1d5 | 382 | 1 | 149.90 | 3 | |
| 4 | 000379cdec625522490c315e70c7a9fb | 153 | 1 | 93.00 | 2 | |
| | Segment | | | | | |
| 0 | At-Risk Customers | | | | | |
| 1 | Big Spenders | | | | | |
| 2 | Big Spenders | | | | | |
| 3 | At-Risk Customers | | | | | |
| 4 | Occasional Buyers | | | | | |

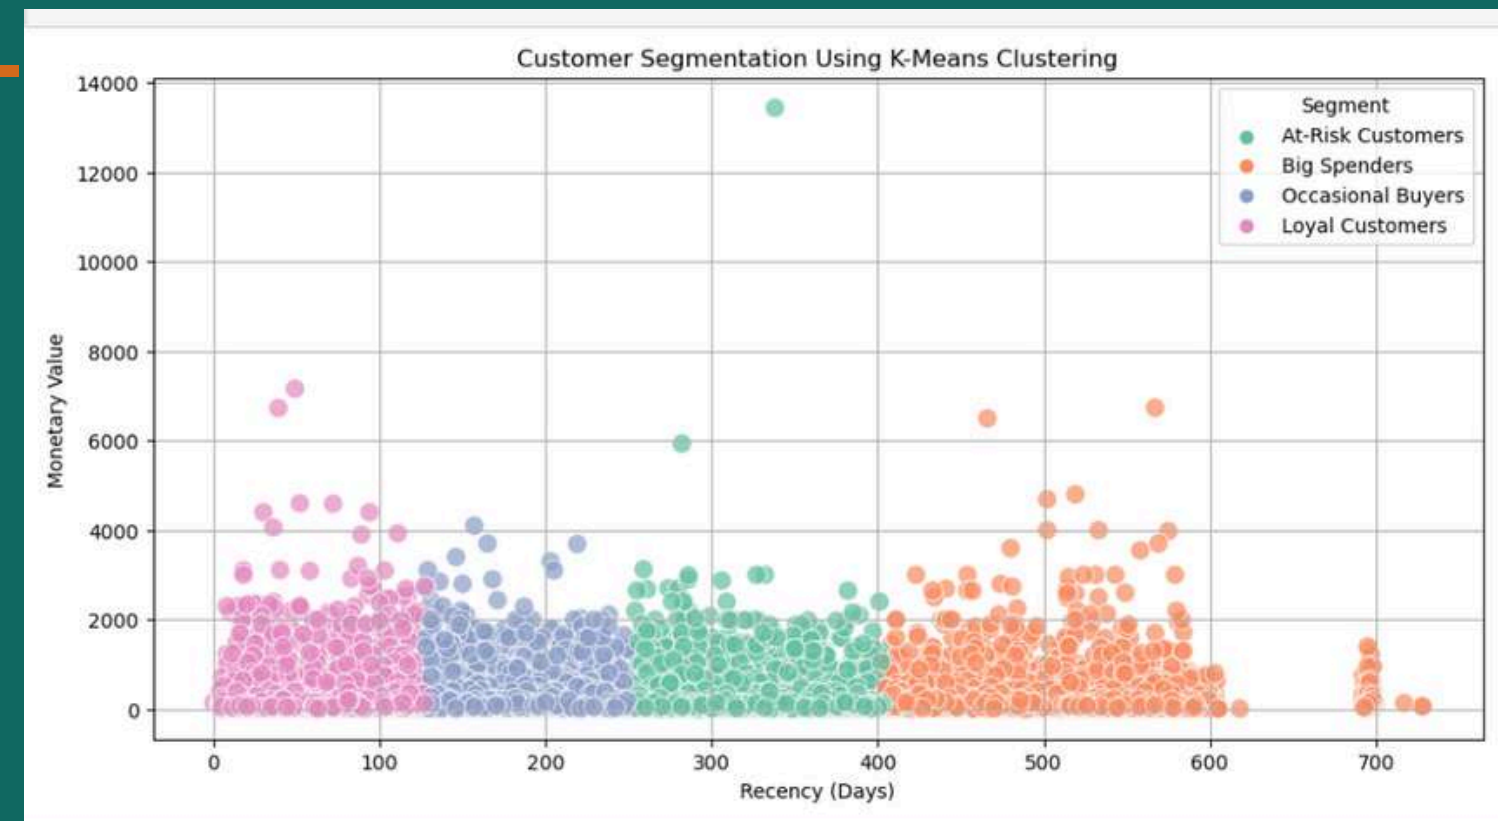
Step 6: Visualize Segments

Visualize customer segments using scatter plots.

```
import seaborn as sns
import matplotlib.pyplot as plt

# Plot the clusters
plt.figure(figsize=(12, 6))
sns.scatterplot(
    x='Recency', y='Monetary', hue='Segment', data=rfm_metrics,
    palette='Set2', s=100, alpha=0.7
)
plt.title("Customer Segmentation Using K-Means Clustering")
plt.xlabel("Recency (Days)")
plt.ylabel("Monetary Value")
plt.legend(title="Segment")
plt.grid(True)
plt.show()
```

Output:-



Step 7: Evaluate the Clustering

Analyze the distribution of customers across segments and their key metrics.

Analyze cluster characteristics

```
cluster_summary = rfm_metrics.groupby('Segment').agg({
    'Recency': 'mean',
    'Frequency': 'mean',
    'Monetary': 'mean',
    'customer_id': 'count'
}).rename(columns={'customer_id': 'Customer Count'})
```

Display the cluster summary

```
print(cluster_summary)
```

Output:-

| Segment | Recency | Frequency | Monetary | Customer Count |
|-------------------|------------|-----------|------------|----------------|
| At-Risk Customers | 319.676714 | 1.149559 | 138.632433 | 25943 |
| Big Spenders | 489.537456 | 1.132596 | 139.143734 | 17821 |
| Loyal Customers | 66.488481 | 1.140131 | 140.219516 | 26304 |
| Occasional Buyers | 188.427967 | 1.141793 | 133.823619 | 28598 |

Optional step: MiniBatch K-Means for Larger Datasets

If the dataset is too large even for K-Means, use MiniBatch K-Means, which is a scalable variation of K-Means

```
from sklearn.cluster import MiniBatchKMeans
```

```
# Perform MiniBatch K-Means clustering
```

```
mini_kmeans = MiniBatchKMeans(n_clusters=n_clusters, random_state=42, batch_size=1000)
```

```
rfm_metrics['Cluster'] = mini_kmeans.fit_predict(rfm_normalized)
```

```
# Display the updated clusters
```

```
print(rfm_metrics[['customer_id', 'Recency', 'Frequency', 'Monetary', 'Cluster']].head())
```

Output:-

```
C:\Users\Dell\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The argument 'n_init' is deprecated in favor of 'n_init=1' or 'n_init='auto''. Please use 'n_init='auto' or 'n_init=1' instead of 'n_init=10'. Set the value of 'n_init' explicitly to suppress the warning.
```

| | customer_id | Recency | Frequency | Monetary | Cluster |
|---|-----------------------------------|---------|-----------|----------|---------|
| 0 | 00012a2ce6f8dcda20d059ce98491703 | 292 | 1 | 89.80 | 3 |
| 1 | 000161a058600d5901f007fab4c27140 | 413 | 1 | 54.90 | 1 |
| 2 | 0001fd6190edaaf884bcacf3d49edf079 | 551 | 1 | 179.99 | 1 |
| 3 | 0002414f95344307404f0ace7a26f1d5 | 382 | 1 | 149.90 | 3 |
| 4 | 000379cdec625522490c315e70c7a9fb | 153 | 1 | 93.00 | 0 |

Q3. Predicting Product Returns:

Build a logistic regression model to predict the likelihood of a product being returned based on features like price, category, and delivery time.

Step 1: Define the Objective

Use the historical data from the datasets (orders.csv, order_items.csv, and products.csv) to predict product returns. Build a logistic regression model to classify whether a product is returned (1) or not (0).

Step 2: Prepare the Data

a) Load the necessary tables Load the orders.csv, order_items.csv, and products.csv tables.

```
import pandas as pd
```

```
# Load datasets
```

```
orders = pd.read_csv("E:\BHU\Data_Analyst\Ecommerce\orders.csv")
```

```
order_items = pd.read_csv("E:\BHU\Data_Analyst\Ecommerce\order_items.csv")
```

```
products = pd.read_csv("E:\BHU\Data_Analyst\Ecommerce\products.csv")
```

```
# Display the first few rows
```

```
print(orders.head())
```

```
print(order_items.head())
```

```
print(products.head())
```

b) Simulate the "product returned" label

Since the dataset does not have explicit return data, we will simulate the return labels:
Assume that orders with a long delivery time and low product price are more likely to be returned. Create a binary column returned (1 for returned, 0 for not returned).

- 1.Calculate delivery time = order_delivered_customer_date - order_approved_at.
- 2.Assign returns based on thresholds for delivery time and product price.

```
import numpy as np
import pandas as pd

# Ensure dates are in datetime format, handling errors
orders['order_approved_at'] = pd.to_datetime(orders['order_approved_at'], errors='coerce')
orders['order_delivered_customer_date'] = pd.to_datetime(orders['order_delivered_customer_date'], errors='coerce')

# Calculate delivery time in days (exclude rows where delivery date is NaT)
orders['delivery_time_days'] = (orders['order_delivered_customer_date'] - orders['order_approved_at']).dt.days

# Merge order_items with orders and products
merged_data = pd.merge(order_items, orders, on='order_id', how='inner')
merged_data = pd.merge(merged_data, products, on='product_id', how='inner')

# Drop rows with missing delivery time or price
merged_data = merged_data.dropna(subset=['delivery_time_days', 'price'])

# Simulate returns (binary target variable)
# Assume products with long delivery time (>10 days) and low price (<20) are more likely to be returned
merged_data['returned'] = np.where(
    (merged_data['delivery_time_days'] > 10) & (merged_data['price'] < 20), 1, 0
)

# Display the prepared dataset
print(merged_data[['product_id', 'price', 'delivery_time_days', 'returned']].head())
```

Output:-

| | product_id | price | delivery_time_days | returned |
|---|----------------------------------|-------|--------------------|----------|
| 0 | 4244733e06e7ecb4970a6e2683c13e61 | 58.9 | 7.0 | 0 |
| 1 | 4244733e06e7ecb4970a6e2683c13e61 | 55.9 | 14.0 | 0 |
| 2 | 4244733e06e7ecb4970a6e2683c13e61 | 64.9 | 17.0 | 0 |
| 3 | 4244733e06e7ecb4970a6e2683c13e61 | 58.9 | 8.0 | 0 |
| 4 | 4244733e06e7ecb4970a6e2683c13e61 | 58.9 | 13.0 | 0 |

Step 3: Feature Engineering

a) Select Features and Target Variable Identify key features for the model:

Features: price, delivery_time_days, and product category. Target: returned.

Convert categorical data (product category) into numerical format using one-hot encoding.

Select relevant columns

```
model_data = merged_data[['price', 'delivery_time_days', 'product category', 'returned']]
```

One-hot encode the product category

```
model_data = pd.get_dummies(model_data, columns=['product category'], drop_first=True)
```

Display the feature matrix

```
print(model_data.head())
```

Output:-

```
price delivery_time_days returned product category_Arts and Crafts \
0 58.9 7.0 0 0
1 55.9 14.0 0 0
2 64.9 17.0 0 0
3 58.9 8.0 0 0
4 58.9 13.0 0 0

product category_Arts and Crafts product category_Bags Accessories \
0 0 0
1 0 0
2 0 0
3 0 0
4 0 0

product category_Blu Ray DVDs product category_CITTE AND UPHACK FURNITURE \
0 0 0
1 0 0
2 0 0
3 0 0
4 0 0

product category_CONSTRUCTION SECURITY TOOLS \
0 0
1 0
2 0
3 0
4 0

product category_Casa Construcáo ... \
0 0 ...
1 0 ...
2 0 ...
3 0 ...
4 0 ...

product category_musical instruments product category_party articles \
0 0 0
1 0 0
2 0 0
3 0 0
4 0 0

product category_perfumery product category_pet Shop \
0 0 0
1 0 0
2 0 0
3 0 0
4 0 0

product category_song product category_sport leisure \
0 0 0
1 0 0
2 0 0
```

b) Split the Data into Training and Testing Sets

Split the dataset into training (80%) and testing (20%) subsets.

```
from sklearn.model_selection import train_test_split
```

```
# Define features (X) and target (y)
```

```
X = model_data.drop('returned', axis=1)
```

```
y = model_data['returned']
```

```
# Split the data
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
print(f"Training data size: {X_train.shape}")
```

```
print(f"Testing data size: {X_test.shape}")
```

Output:-

```
Training data size: (88156, 74)
```

```
Testing data size: (22040, 74)
```

Step 4: Train the Logistic Regression Model

Train the logistic regression model on the training data.

```
from sklearn.linear_model import LogisticRegression
```

```
# Initialize the logistic regression model
```

```
logistic_model = LogisticRegression(max_iter=1000, random_state=42)
```

```
# Train the model
```

```
logistic_model.fit(X_train, y_train)
```

```
# Display model coefficients
```

```
feature_names = X.columns
```

```
coefficients = logistic_model.coef_[0]
```

```
for name, coef in zip(feature_names, coefficients):
```

```
    print(f"{name}: {coef}")
```

Output:-

```
price: -0.28798894517165935
delivery_time_days: 0.16137177776462228
product category_Art: -0.6295351166073053
product category_Arts and Crafts: -0.39923378451169445
product category_Bags Accessories: 0.7007639124092929
product category_Blu Ray DVDs: 1.0812552655603513
product category_CITTE AND UPHACK FURNITURE: -0.34793037121252346
product category_CONSTRUCTION SECURITY TOOLS: 0.5465364866074903
product category_Casa Construciao: -0.17285817416621965
product category_Christmas articles: -0.7136892389216409
product category_Construction Tools Construction: 0.2756021343020453
product category_Construction Tools Garden: 0.8451076661127886
product category_Construction Tools Illumination: -0.46721922178048864
product category_Construction Tools Tools: -0.15924555474103513
product category_Cool Stuff: -0.060018194442146164
product category_Drink foods: 0.3369075474561295
product category_ELECTRICES 2: 0.4853652868898182
product category_Fashion Bags and Accessories: -0.2959207995896664
product category_Fashion Calcados: -0.4639501217974638
product category_Fashion Children's Clothing: -0.00015298246259712084
product category_Fashion Men's Clothing: -0.6309511553629373
product category_Fashion Sport: -0.04416706119204735
product category_Fashion Underwear and Beach Fashion: -2.343233697900357
product category_Fashion Women's Clothing: 0.9503730533496866
product category_Furniture: -1.638050068724108
product category_Furniture Decoration: 0.31625344842123637
product category_Furniture Kitchen Service Area Dinner and Garden: 0.7585288585340342
product category_Furniture office: -0.06845379302726698
product category_Games consoles: 1.3153238755219034
```


Step 5: Evaluate the Model

a) Make Predictions Predict on the testing data.

Make predictions

```
y_pred = logistic_model.predict(X_test)
```

```
y_pred_prob = logistic_model.predict_proba(X_test)[:, 1] # Probability of return
```

b) Evaluate the Model's Performance

Use accuracy, precision, recall, and F1-score to evaluate performance.

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
```

Calculate evaluation metrics

```
accuracy = accuracy_score(y_test, y_pred)
```

```
precision = precision_score(y_test, y_pred)
```

```
recall = recall_score(y_test, y_pred)
```

```
f1 = f1_score(y_test, y_pred)
```

```
roc_auc = roc_auc_score(y_test, y_pred_prob)
```

Print metrics

```
print(f"Accuracy: {accuracy:.2f}")
```

```
print(f"Precision: {precision:.2f}")
```

```
print(f"Recall: {recall:.2f}")
```

```
print(f"F1-Score: {f1:.2f}")
```

```
print(f"ROC-AUC: {roc_auc:.2f}")
```

Output:-

```
Accuracy: 0.98  
Precision: 0.67  
Recall: 0.49  
F1-Score: 0.57  
ROC-AUC: 0.99
```

Step 6: Interpret Results

1. Analyze which factors contribute the most to predicting returns:
2. Feature Importance: Use the coefficients from the logistic regression model.
3. Customer Patterns: Examine the most common conditions leading to product returns.

Step 7: Optional - Visualize Results

a) Plot the ROC Curve Visualize the model's ability to distinguish between returned and non-returned products.

```
from sklearn.metrics import roc_curve
```

```
# Compute ROC curve
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
```

```
# Plot ROC curve
```

```
plt.figure(figsize=(8, 6))
```

```
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {roc_auc:.2f})", color='blue')
```

```
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
```

```
plt.xlabel("False Positive Rate")
```

```
plt.ylabel("True Positive Rate")
```

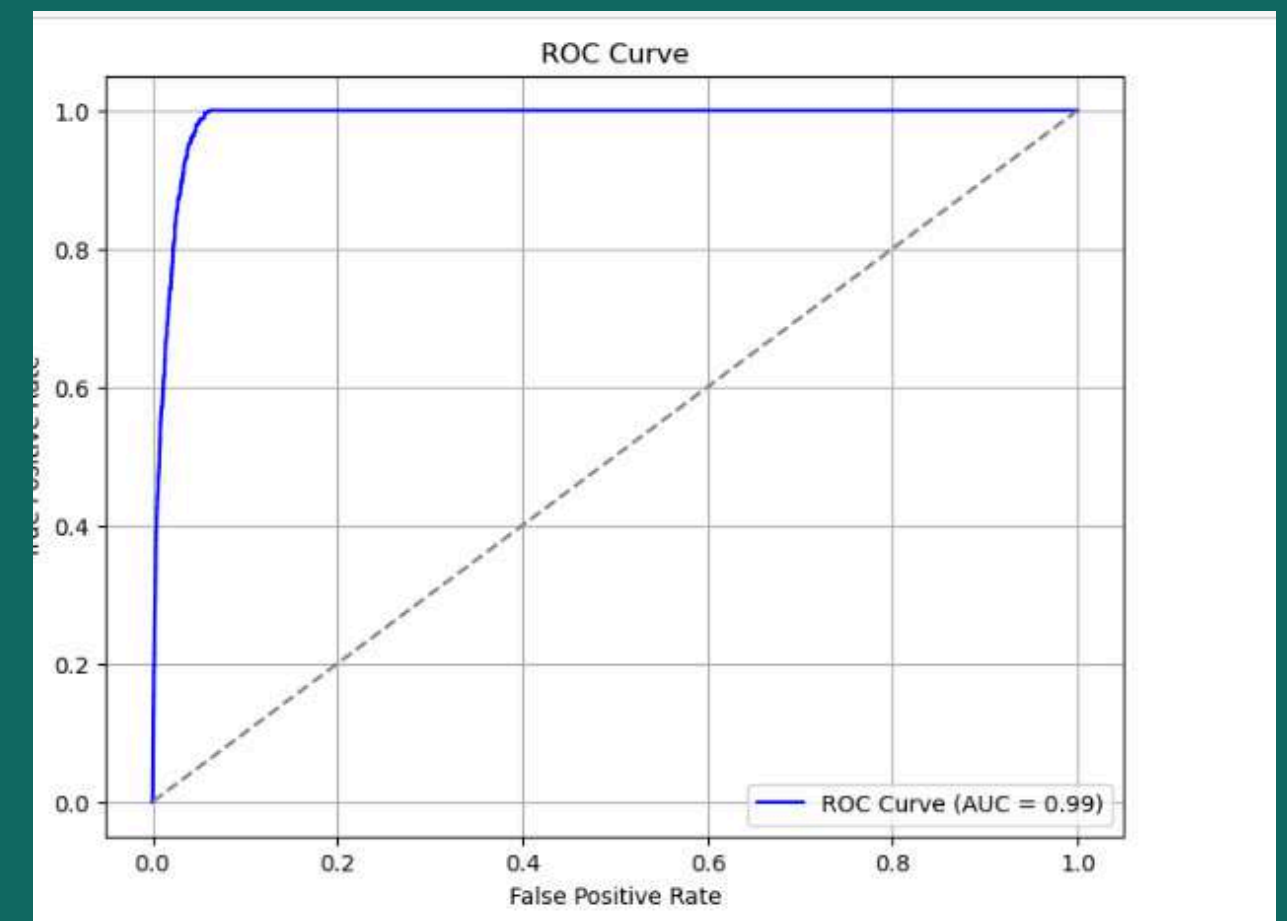
```
plt.title("ROC Curve")
```

```
plt.legend()
```

```
plt.grid()
```

```
plt.show()
```

Output:-



Q4. Cross-Sell/Upsell Opportunities:

Use correlation analysis and simple rules-based approaches to identify frequently purchased product combinations.

Provide recommendations for cross-selling or bundling products.

Steps

- Merge Necessary Datasets We need to combine order_items with products to associate product details with each order.
- Identify Frequently Purchased Product Combinations Group by order_id and analyze which products are purchased together. Use a co-occurrence matrix or market basket analysis (e.g., Apriori algorithm) to identify frequent combinations.
- Correlation Analysis Compute a co-purchase correlation matrix to quantify relationships between product pairs.
- Provide Recommendations Based on the analysis, suggest product bundles or cross-selling strategies.

code:-

```
from scipy.sparse import coo_matrix  
import numpy as np  
from sklearn.metrics.pairwise import cosine_similarity  
import pandas as pd
```

Step 0: Merge order_items with product details

```
merged_items = pd.merge(order_items, products, on='product_id', how='inner')
```

Clean and inspect the merged data

```
merged_items = merged_items.dropna(subset=['order_id', 'product_id']) # Remove rows with missing order_id or product_id
```

```
merged_items['order_id'] = merged_items['order_id'].astype(str) # Ensure order_id is a string
```

```
merged_items['product_id'] = merged_items['product_id'].astype(str) # Ensure product_id is a string
```

Step 1: Extract unique order_ids and product_ids

```
order_ids = merged_items['order_id'].unique()
```

```
product_ids = merged_items['product_id'].unique()
```

Step 2: Create the binary sparse matrix

Map orders and products to unique integer indices

```
order_index = {order_id: idx for idx, order_id in enumerate(order_ids)}
```

```
product_index = {product_id: idx for idx, product_id in enumerate(product_ids)}
```

Map rows and columns for the sparse matrix

```
rows = merged_items['order_id'].map(order_index)
```

```
cols = merged_items['product_id'].map(product_index)
```

```
data = np.ones(len(merged_items), dtype=int)
```

Build a sparse binary matrix

```
sparse_matrix = coo_matrix((data, (rows, cols)), shape=(len(order_ids), len(product_ids)))
```

Step 3: Compute the co-purchase correlation matrix

Transpose and convert the sparse matrix to CSR format for efficient row-based operations

```
dense_matrix = sparse_matrix.T.tocsr()
```

Compute pairwise product correlations using cosine similarity

```
correlation_matrix = cosine_similarity(dense_matrix, dense_matrix)
```


Step 4: Identify strong correlations for recommendations

Flatten and filter strong correlations

```
product_ids_array = np.array(product_ids) # Convert product IDs to a numpy array
correlated_pairs = [] # List to store correlated pairs
```

```
threshold = 0.6 # Define the threshold for strong correlations
for i in range(correlation_matrix.shape[0]):
    for j in range(i + 1, correlation_matrix.shape[1]): # Avoid self-correlations
        if correlation_matrix[i, j] > threshold:
            correlated_pairs.append(
                (product_ids_array[i], product_ids_array[j], correlation_matrix[i, j])
            )
```

Sort by correlation in descending order

```
correlated_pairs = sorted(correlated_pairs, key=lambda x: x[2], reverse=True)
```

Step 5: Display cross-sell/upsell recommendations

```
print("Cross-Sell/Upsell Opportunities:")
for pair in correlated_pairs[:10]: # Display top 10 recommendations
    print(f"Recommend bundling Product {pair[0]} with Product {pair[1]} (Correlation: {pair[2]:.2f})")
```

Provide actionable recommendations

```
print("\nActionable Recommendations:")
for pair in correlated_pairs[:5]: # Focus on the top 5 for actionable insights
    print(
        f"Consider bundling or cross-selling Product {pair[0]} with Product {pair[1]}."
        f"These products have a strong purchase correlation of {pair[2]:.2f}."
    )
```

Output:-

```
Cross-Sell/Upsell Opportunities:
Recommend bundling Product d143bf43abb18593fa8ed20cc990ae84 with Product 55939df5d8d2b853fbc532bf8a00dc32 (Correlation: 1.00)
Recommend bundling Product 84bc1370758fa98d08afd0180818217a with Product 18692fdc7adb29c885d8cb84eb68a97e (Correlation: 1.00)
Recommend bundling Product 5edf955f0c8417e3fea7288360c2e22d with Product 188c9cb0aff9e485e3f84752731c633a (Correlation: 1.00)
Recommend bundling Product dd12801e4d1919e84bf6f90167f22fb5 with Product bc0d2d816d425e6ce3e80d8b9887d324 (Correlation: 1.00)
Recommend bundling Product 683487a7a9ad71b5776c9b9bcad37eed with Product 30b159b653b87754d7a38446c7ef4669 (Correlation: 1.00)
Recommend bundling Product ba2fafdf0c3e7418202575a48d00bd15 with Product e87093a08e761c9fbf7dd461227f66c1 (Correlation: 1.00)
Recommend bundling Product 62c2b9bd44500d0305b1e50e2c9bd34d with Product cb790fa02cc6097dcfbf8f9f05904a005 (Correlation: 1.00)
Recommend bundling Product 095fc85516c1210e6dbed95ed8041ee0 with Product b7f9e962dba31973edd059ada998d699 (Correlation: 1.00)
Recommend bundling Product fad4c004f68304a86535ea71c79cf4aa4 with Product 4374f5c6e8ab35133a8dbc5517ec0388 (Correlation: 1.00)
Recommend bundling Product d961667356e6dffe33c0d0eb4c327cc with Product 93c285dad77a797cf2193ead6afbe107 (Correlation: 1.00)

Actionable Recommendations:
Consider bundling or cross-selling Product d143bf43abb18593fa8ed20cc990ae84 with Product 55939df5d8d2b853fbc532bf8a00dc32. These products have a strong purchase correlation of 1.00.
Consider bundling or cross-selling Product 84bc1370758fa98d08afd0180818217a with Product 18692fdc7adb29c885d8cb84eb68a97e. These products have a strong purchase correlation of 1.00.
Consider bundling or cross-selling Product 5edf955f0c8417e3fea7288360c2e22d with Product 188c9cb0aff9e485e3f84752731c633a. These products have a strong purchase correlation of 1.00.
Consider bundling or cross-selling Product dd12801e4d1919e84bf6f90167f22fb5 with Product bc0d2d816d425e6ce3e80d8b9887d324. These products have a strong purchase correlation of 1.00.
Consider bundling or cross-selling Product 683487a7a9ad71b5776c9b9bcad37eed with Product 30b159b653b87754d7a38446c7ef4669. These products have a strong purchase correlation of 1.00.
```

Resource Page





Target
Insight



THANK YOU

For watching this presentation

K K Sahani

+91 72557964940

krishnakumarsahani.bhumca22@gmail.com

<https://www.linkedin.com/in/krishna-kumar-sahani-79867225a/>

chandigarh, India